Supplement S2: Extended methods and complete historical results

Anonymous authors

Interpretation update. This extended methods record preserves historical comparisons. Supplement S1 and the revised main paper add post-review statistical-only pooling controls and CAL-selection sensitivity. Those results do not establish an LLM-specific advantage; they supersede broader readings of the earlier fixed-mixture gains.

Abstract

Does a persona improve prediction of an individual's unseen questionnaire answers beyond statistical use of the same observed answers? In an 84-template, two-model experiment, the planned comparison did not establish superiority of the MAP-Elites-style optimizer. On a further 600 people from the same inventory, a frozen Granite persona–ridge mixture improved normalized ranked probability score (RPS) by 0.00125–0.00184 against two neutral mixtures. On reanalysis of these same respondents, the result depended on the statistical comparator's training budget. The fixed hybrid led at 80 training people, with additional historical supervision for the hybrid. At every tested budget from 320 onward, statistical predictors led in both representations. At 8,320 people the hybrid's excess RPS was 0.021679 for raw answers and 0.018311 for categorical profiles. With that training cohort fixed, the raw-answer decoder matched the 60-answer hybrid at an interpolated 11.44 known answers (respondent-bootstrap 95% interval [10.46, 12.38], conditional on three item orders). An auxiliary analysis on an external survey aligned each respondent's forecast with the actual product and price. On 200 calibration respondents, mean log loss was 0.5192 for the statistical decoder versus 0.6373 for product-specific priors. On 658 held-out respondents, only the two CAL-selected Granite mixtures improved on the full-data statistical control after multiplicity adjustment. The normalized reductions were small (0.00404 and 0.00384). Transferring the persona instructions improved standalone Granite and worsened Qwen. The direct post hoc comparison of CAL-selected Granite mixtures was near zero, with its paired interval spanning zero.

1 Introduction

An agent that answers a questionnaire in character is useful as an individual predictor only if its answers forecast a particular unseen person better than simpler uses of the same human information. Population-level resemblance, plausible prose, and a high score on people used for prompt selection do not establish that property. This is a practical issue for survey simulation, where the relevant unit of analysis is the individual and acquiring respondent-level self-reports entails time and compensation costs Park et al. [2026]. Prior work on personality-conditioned agents evaluates induced behavioral style through changes in task selection Newsham and Prince [2025]; our target is the human associated with a profile.

We study a reusable *persona template*: a prompt optimized on development respondents, then instantiated from a new respondent's 60 observed answers to forecast 60 different answers. The template cannot use that person's target answers, target-derived cluster, or evaluation outcomes. We ask whether proper-score search improves predictions under this restriction and whether any improvement exceeds budget-matched search and statistical predictors with access to the same observed items.

We separate three questions that require different controls. First, does a respondent's own profile predict their answers better than a different respondent's profile? Second, with the history fixed, does searching over persona wording improve the forecast over a neutral instruction? Third, what would a statistical predictor achieve with the same observed answers and the available labeled training respondents? A successful profile match answers the first question; it need not imply either of the other two.

The original planned optimizer comparison was null. A narrower result did survive transfer to a third respondent panel from the same instrument: the frozen Granite persona–ridge recipe, with its selected calibration and mixture weight, beat two neutral recipes when the ridge component was fitted to 80 people. We then compared that fixed recipe with statistical models chosen by cross-validation at six training budgets on the same, previously evaluated panel (Figure 1). The hybrid leads at 80. At 160, both point estimates favor the hybrid; only the raw-answer interval includes zero. The statistical predictors lead at every tested budget from 320 onward. The hybrid also used historical human labels for search, selection, calibration and population priors. Its 80-person ridge fit therefore understates its total supervision.

The original search, fitness-objective and profile-matching decisions are reported under their original questions and multiplicity families below. For the separate Twin-2K-500 instrument, an outcome-blind source audit recovered each respondent's actual product and price. We evaluate statistical predictors of these offers on 200 calibration respondents and retain the earlier history-only analysis as a diagnostic. On the registered test split, CAL-

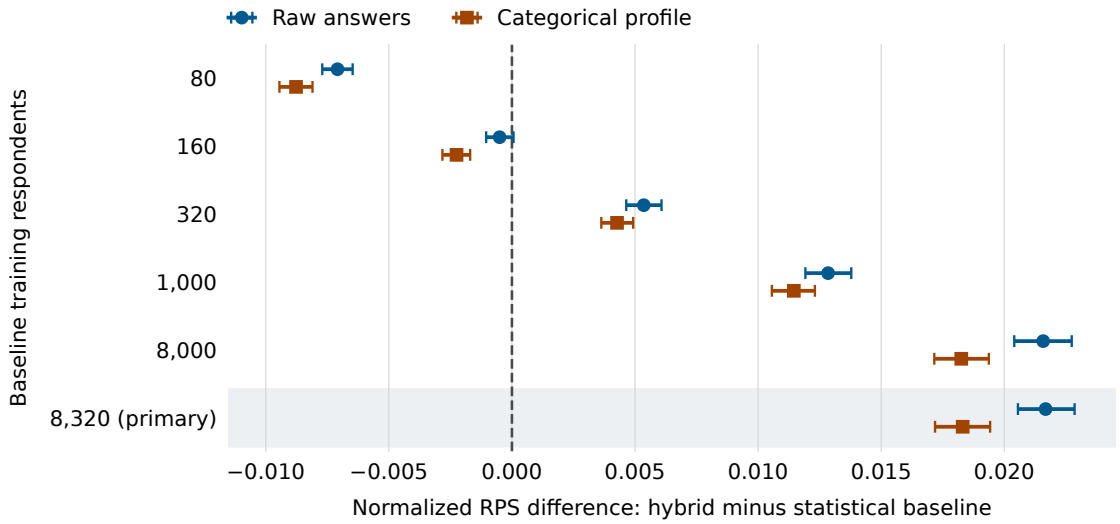


Figure 1: Effect of the fixed hybrid relative to a within-budget selected statistical predictor. Points are differences in normalized ranked probability score; smaller scores are better. Intervals resample the same 600 respondents, conditional on the fitted models. Both representations at 8,320 respondents form the two declared primary contrasts. The other budgets are secondary, and the E3 cohort had already been evaluated.

selected Granite mixtures improved modestly against the full-data control, with intervals spanning zero for the corresponding Qwen mixtures. The public Wave-4 labels and unknown pretraining exposure limit the strength of the transfer evidence.

Companion supplement map. Supplement S1 contains the latest post-review controls and qualifies the earlier conditional interpretation. This S2 preserves extended historical methods and results. S3 contains the external source protocol and B1/R2 tables. S4 preserves historical Tables S1–S16 and the E1–E5 analyses with their original table numbering; its pagination has changed during neutral reformatting.

2 Related work

Personality and agent behavior. Newsham and Prince Newsham and Prince [2025] study how induced OCEAN traits alter task selection by LLM agents. Xu et al. Xu et al. [2023] condition narrative dialogue on personality representations and evaluate it with automatic measures and human judgments. These studies ask whether an agent expresses a supplied personality. Our outcome is a probability forecast for a particular held-out human, so trait expression alone cannot validate it. Jahan and Mell Jahan and Mell [2024] relate personality and self-reports to behavior in human–agent negotiation. That behavioral criterion is distinct from our questionnaire target: our results do not establish transfer to negotiation or other interactive decisions.

Individual prediction and prompt search. Persona-based simulation Hu and Collier [2024], Park et al. [2026] and learned models of human cognition Binz et al. [2025] motivate evaluating human outcomes directly. We isolate a narrower intervention: changing reusable persona instructions while fixing the observed items and readout model. Hu and Collier compare persona prediction with supervised linear regression. We extend that comparison to fitting budget, exact rendered features and search-objective choice. Recent work tests temporal prediction Jia et al. [2026] and attribute-dependent population alignment Fröhling et al. [2026]; we isolate prompt optimization within one inventory. Aher et al. Aher et al. [2023] instead evaluate replication of human-subject study patterns; our target is each held-out respondent's disjoint questionnaire answers. Argyle et al. Argyle et al. [2023] study demographic conditioning of synthetic survey samples. Bisbee et al. Bisbee et al. [2024] find that similar survey means can coexist with distorted variation and associations; Tjuatja

et al. Tjuatja et al. [2024] show that survey response biases often differ between LLMs and humans. We therefore test matched individual forecasts, without treating predictive gains as validation of human-like behavior. Our search families draw on EvoPrompt Guo et al. [2024], PromptBreeder Fernando et al. [2024], OPRO Yang et al. [2024], ProTeGi Pryzant et al. [2023], and GEPA Agrawal et al. [2026]. Calibration-aware selection also appears in Reflective Prompt Tuning Bayat et al. [2026]; we test objective effects under fixed information and budget.

Unlike the adversarial scenario search in MADRID Mouret and Clune [2015], Samvelyan et al. [2024], we optimize individual response forecasts.

3 Task and method

Figure 2 separates the per-person forecast from the development, freezing and evaluation stages.

3.1 Respondent forecasting

Let x_i be a new respondent's answers to the fixed 60 input items and $y_{ij} \in \{1, \ldots, 5\}$ the answer to target item j from the disjoint 60-item set. A frozen template g constructs a persona $g(x_i)$, and a pinned readout model returns an ordinal distribution $p_{ij}(k)$. We evaluate the normalized ranked probability score (RPS), a proper scoring rule for ordinal forecasts Gneiting and Raftery [2007]

$$\mathrm{RPS}(p, y) = \frac{1}{K-1} \sum_{k=1}^{K-1} \left(\sum_{h=1}^{k} p(h) - \mathbf{1}\{y \le k\} \right)^2, \quad K = 5. \tag{1}$$

Lower is better. We average the 60 target-item losses within each respondent before averaging respondents; a search seed is a repeated optimization run, not a new person. Our primary estimand is the respondent-population mean paired change in calibrated RPS, averaged over the three frozen search runs. Inference over fresh optimizer runs is a separate target. We retain five-category probability estimates and digit-token mass so that losses and readout coverage can be checked separately.

3.2 Information available to the persona

The IPIP-NEO-120 inventory measures 30 facets of five personality traits Johnson [2014]. The renderer converts the observed input items into facet means on a 0–100 scale, averages facets within each trait, and maps these scores to fixed verbal intensity modifiers. A template supplies a role definition, five trait formulations, 30 facet formulations, and an internal-reflection instruction. Search edits these text fields; the intensity mapping and item split are fixed. The neutral base uses the same renderer with standard construct names and instructions to answer from the supplied profile. No target-derived cluster selects a template.

Thus the LLM receives a compressed verbal profile, not a transcript of all 60 answers. All LLM search arms share this representation. The statistical reference receives the same observed items as one-hot encoded answers. This matches available human history but does not match representation: any LLM–ridge gap combines the effects of compression, readout, and learning. We therefore add a post-hoc predictor fitted to the exact categorical profile to determine whether compression alone can explain that gap. The comparison still differs in fitted predictor and training objective, so it cannot isolate an intrinsic limitation of the LLM architecture.

3.3 Readout and calibration

Each target is presented with a five-point response scale using the model's chat template and the assistant prefill `My answer is`. The readout requests one token at temperature zero with top-20 log probabilities. It sums the returned probabilities for answer digits 1–5 and normalizes their mass; if their total mass is zero, the implemented fallback is uniform. These are estimates derived from returned digit probabilities, not exact full-vocabulary probabilities. Reverse-keyed items are mapped onto the corpus's recoded scale before scoring. Digit mass and the fraction below 0.5 are retained separately.

Let Q_{ij} denote the model's four-threshold CDF, F_{0j} the empirical CDF from training respondents, and $\bar{Q}_j^C$ the mean model CDF on the calibration panel. The evaluated forecast is

$$\widetilde{Q}_{ij} = \Pi_{\mathcal{C}} \left[F_{0j} + \alpha \left(Q_{ij} - \bar{Q}_j^C \right) \right], \tag{2}$$

where $\Pi_{\mathcal{C}}$ minimizes squared error over CDFs bounded by 0 and 1. For each selected template and the base, α minimizes calibration-panel RPS over $\{0, .1, \ldots, 1, 1.2, 1.4, 1.6, 1.8, 2\}$. The resulting α and centering vector are then fixed on evaluation respondents. Wrong-persona forecasts use the evolved template's calibration parameters unchanged. The training labels therefore contribute to LLM calibration as well as to the statistical references; this is a target-supervised evaluation.

The main search objective is *raw* RPS. In calibrated-fitness controls, α and the centering vector are fitted on one fixed half of the search panel and scored on the other half, then swapped. These two folds are shared across candidates and seeds. The final calibration panel is not used during search. All variants still use raw RPS for shortlist selection. The legacy GA control minimizes the negative expected proximity $-\left[1 - \sum_{k=1}^{5} p(k)|k - y|/4\right]$, rather than RPS.

3.4 Persona search

The two readout families are Qwen3.6-35B-A3B and Granite-4.2-30B; a pinned Qwen3.8-27B mutator proposes edits. The eleven methods are GA, DE, PromptBreeder, MAP-Elites-style search, GEPA-style search, offi-

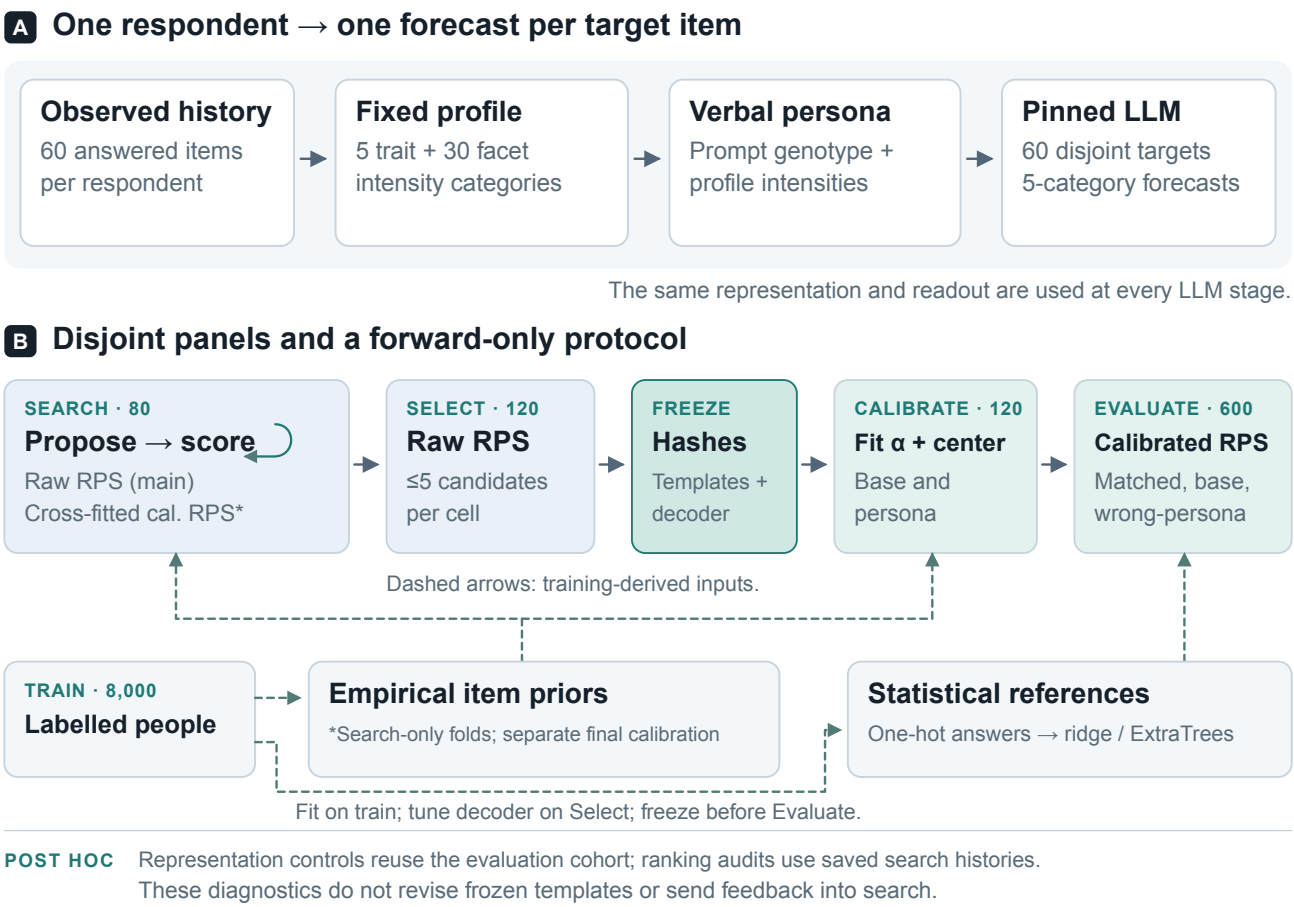


Figure 2: Persona forecasting and frozen evaluation. (A) Sixty observed answers become 35 categorical profile values; searched wording and a pinned readout produce ordinal forecasts for 60 disjoint items. (B) Respondent panels are disjoint. Search uses raw or cross-fitted calibrated RPS, but every shortlist is selected by raw RPS. Templates and decoder choices are frozen before final calibration and evaluation. Dashed arrows show training-derived inputs. Post-hoc diagnostics do not revise templates or feed back into search. The legacy GA fitness control is omitted for readability.

cial GEPA, native GigaEvo, OPRO, ProTeGi, random search, and best-of-budget. GA and DE adapt EvoPrompt's evolutionary operators to the common JSON schema. PromptBreeder evolves mutation instructions; OPRO uses scored history; ProTeGi uses textual critiques. Random search independently mutates the seed template, while best-of-budget samples initial candidates without an evolutionary trajectory. Local adapters share the task-specific genotype, feedback, and evaluator; they are not replications of the original papers' benchmark results. Official GEPA v0.1.4 and native GigaEvo v2.20.1 are pinned separately from their local style-based counterparts.

The MAP-Elites-style adapter uses two islands and an archive indexed by prompt length and predicted-response standard deviation. It starts with four candidates per island and mutates candidates until the forecast budget is exhausted, with migration every 12 iterations and a six-observation insight window. These are fixed search settings, not evidence that archive diversity causes an improvement. The GA also stops on the forecast budget; its settings are population eight, two elites, tournament size three, crossover probability 0.8, and mutation probability 0.3. The common genotype-proposal settings are temperature 0.7 and a 2,500-token output cap; auxiliary operator requests can use different settings.

One full-panel forecast equivalent is $80 \times 60 = 4{,}800$ respondent–item forecasts. Each cell is capped at 80 equivalents, or 384,000 executed candidate forecasts; fractional evaluations consume their actual share. Shortlist selection has a separate cap of five candidates on 120 people. Failed proposals and their available token costs are retained separately; a proposal is not credited with forecasts that did not occur. Equal forecast caps do not imply equal tokens, mutator calls, or wall time.

Table 1: Fixed respondent roles. IDs are disjoint across rows; the 60 observed items and 60 target items are disjoint within every row.

Split	People	Permitted use
Train	8,000	Priors and statistical fitting
Search	80	Candidate fitness
Select	120	Prompt and decoder selection
Calibrate	120	Final LLM calibration
Evaluate	600	Frozen comparison

4 Experimental design

4.1 Splits, freezing, and controls

We use the complete-case IPIP-NEO-120 corpus with split seed 20261004. A source audit matches all 410,168 retained records to the original archive by respondent ID and item values (Supplement S4). Table 1 gives respondent roles. Previously used historical IDs, saved response fingerprints, and baseline-panel IDs are excluded according to the available exposure inventory. This reduces within-program exposure but does not establish that the public corpus was absent from model pretraining or earlier methodological development.

The design has 66 main cells (11 methods, two models, seeds 1–3) and 18 additional cells (GA with calibrated or legacy fitness and MAP-Elites-style search with calibrated fitness). The fixed manifest names raw-fitness MAP-Elites-style search as primary. Within each model, the lowest mean selection-panel raw RPS across seeds chooses the strongest non-primary search comparator, excluding objective variants. This yields OPRO for Qwen3.6 and native GigaEvo for Granite. The 84 selected templates and the statistical decoder configuration are hash-bound before final evaluation. These are internal protocol and freeze records, not an assertion of external preregistration.

Every final cell evaluates the frozen template and neutral base on the same 600 people. Base forecasts are recomputed in each cell; the base score in a paired comparison is the corresponding cell's score, not a pooled reference from another run. A common derangement assigns another person's profile to each respondent, with no fixed points. It retains the target item and model, measuring whether matching the profile to the respondent matters.

The statistical controls are an empirical item prior, a multioutput ridge ordinal-CDF predictor, and ExtraTrees. Both learned references use one-hot encoding of the 60 observed answers and are fitted on the 8,000 training people. Ridge selects its penalty from $\{1, 10, 100, 1000\}$; ExtraTrees uses 64 trees and chooses minimum leaf size from $\{10, 30\}$. Selection-panel RPS chooses hyperparameters and the strongest family (ridge). Unlike LLM readouts, these references use no calibration targets; the comparison matches input history, not total supervision. Evaluation people fit neither predictor. Predicted CDFs are projected onto valid ordinal distributions before scoring.

Retained-information curves use ridge, three frozen nested input orders, and budgets $m \in \{0, 1, 2, 5, 10, 20, 30, 40, 60\}$. Ridge is tuned separately at each positive budget. Monotonic regression and interpolation translate the persona loss into an answer equivalent. This is information retained under a specified decoder, not the number of questionnaire answers literally saved.

4.2 Uncertainty and decision rule

The unit of inference is a respondent from the filtered corpus, with the same people paired across methods and seeds. Estimates condition on archived readouts. For losses L_{is}^A and L_{is}^B, the estimated contrast is $\widehat{\Delta} = N^{-1} \sum_i [3^{-1} \sum_{s=1}^3 (L_{is}^A - L_{is}^B)]$. Raw and calibrated runs are paired by initialization seed index and share respondent panels and calibration folds. Their adaptive search paths differ; this pairing does not ensure identical mutator samples or serving execution. We use 5,000 crossed bootstrap draws resampling both respondents and the three search seeds, and 5,000 two-sided respondent-level sign flips after averaging the three seeds. The sign-flip test conditions on the three frozen runs; the crossed interval additionally describes sensitivity to resampling those observed runs, not validated uncertainty for new optimizer runs. Thus a small conditional p value can accompany a crossed interval containing zero. Holm correction covers all 110 comparisons in the analysis: 84 arm-versus-base/decoder/wrong-persona contrasts and 26 primary-versus-other-arm contrasts across the two models. Displaying a subset does not change that correction family. At this simulation resolution, several adjusted values equal approximately .022; they are not precise estimates of identical underlying probabilities.

We require a crossed interval excluding zero in the favorable direction and an adjusted $p < .05$ before describing a contrast as an improvement. These are different uncertainty calculations, not an interval and test obtained by inverting the same procedure. A benefit relative to base is distinct from practical superiority to the strongest same-history statistical reference. A missing final prediction fails a cell instead of dropping a difficult respondent. Hash checks enforce a common panel, manifest, selected templates, and decoder before combining results.

4.3 Additional-respondent fixed-template transfer

After inspecting the original cohort, we fixed a separate test of 24 historical GA and MAP-Elites-style winners (two models, raw or calibrated fitness, three old seeds) on 600 additional people from the same IPIP120 corpus. The new IDs were sampled from 36,119 complete eligible cases after excluding the reachable inventoried historical exposure union; unknown lifetime and pretraining exposure remain possible. Original Train8000 priors, Calibrate120 weights and CDF centers, model snapshots, selected templates,

and ridge specification were retained. No new prompt search or selector intervention was run. The internally frozen primary family comprises three Granite contrasts: calibrated versus raw MAP-Elites-style winners, calibrated MAP-Elites-style versus its neutral base, and calibrated GA versus its neutral base. We apply the same favorable crossed-interval and Holm-$p < .05$ rule, correcting the three sign-flip tests together. Qwen, wrong-profile, same-history ridge, and other arm contrasts are reported as diagnostics. The archival Part B protocol and full results separate this test from the original 110-comparison family.

4.4 Exploratory same-shortlist selector study

After inspecting the original evaluation, we re-scored the same five archived candidates in each of 12 GA searches (two models, raw/calibrated search fitness, three old seeds) on the original Select120. Fresh forecasts gave both a raw-RPS and a two-fold cross-fitted calibrated-RPS selector per pool; all winner identities were frozen before re-evaluating them, with fresh neutral bases, on the *already exposed* original Evaluate600. Each winner refits calibration on the original Calibrate120. Four direct selector effects, four search-objective effects and two interactions form a separate ten-test, two-sided Holm family. Crossed intervals resample respondents and the three observed seeds; sign flips condition on those archived runs. This study neither supplies new people nor tests new search trajectories or E1's additional respondents (Supplement S4).

4.5 Frozen Granite-80 mixture transfer

After inspecting the original, additional-respondent, and selector outcomes, we internally froze a narrower test on a third 600-person panel from the same IPIP120 corpus. Its IDs exclude every original split, the first additional panel, and the reachable historical exposure inventory. Only IDs and item missingness determined eligibility; target answers were opened after the old recipe and panel were frozen. We reused the three historical Granite calibrated-GA prompts, their neutral counterparts, old calibration and population prior, and six ordinary and six prior-anchored ridge fits for categorical/raw representations and three disjoint 80-person training subsets. Old Calibrate120 was split into 60 calibration people and 60 people for independent mixture-weight selection. Ridge penalties used old Select120. The 80 people are a paired *ridge fitting* budget, not total target-label supervision: search, selection, prior, calibration and each new person's 60 observed answers also contribute information.

For each representation, we compare the optimized persona plus ordinary ridge with neutral persona plus (i) the strongest prior-anchored ridge and (ii) the same ordinary ridge. The neutral arms retain their own previously selected weights. Thus (ii) holds ridge coefficients fixed but tests the optimized prompt-plus-weight recipe, not a prompt-only intervention. The four comparator-minus-hybrid RPS contrasts form one two-sided Holm family. We average nine fixed per-person *loss* vectors (three old search seeds by three fixed ridge subsets) before paired respondent inference. A representation passes only if both gains are at least 0.001 RPS, both paired 95% bootstrap lower bounds exceed zero and both adjusted respondent sign-flip p values are below .05. Intervals condition on those old fits and this panel; they do not quantify new optimizer-run variability. The protocol, full controls and source lineage are in Supplement S4.

4.6 Retrospective same-panel capacity comparison

After E3 was exposed, we specified a separate B1 question: how does its *fixed hybrid recipe* compare with statistical predictors given the same 60 observed test-person items as the number of fitted respondent mappings increases? We fit prior, ridge and nonlinear families from either the raw item answers or the exact rendered categorical profile. At each nested training budget (80, 160, 320, 1,000, 8,000, 8,320 distinct people), three fixed subset seeds receive their own three-fold development-only cross-validation; the best family and hyperparameters are frozen before E3 scoring. We verify all 18 fit receipts, 108 method rows, source hashes, person/item IDs and labels. The B1 comparison uses the mean of three fixed baseline loss vectors and the mean of nine correlated, historical E3 hybrid loss vectors per respondent; it does not average forecast probabilities or resample those combinations as independent searches. The two n=8,320 representation contrasts form the primary Holm family; twelve budget by representation contrasts form a separate secondary family. Conditional respondent bootstrap intervals and sign flips follow this fixed-run estimand. A gain gate requires hybrid-minus-baseline RPS at most -0.001, its CI upper bound below -0.001, and Holm-adjusted $p < .05$. All budgets are retained regardless of outcome. This is a retrospective comparison on an already exposed cohort. The hybrid's historical search, prior, selection, calibration and pretrained-model information is not equalized by the baseline fit count.

4.7 External self-report protocol and stimulus validity

A separate Twin-2K-500 protocol Toubia et al. [2025] defines 49 Wave-4 targets: 40 hypothetical priced purchase choices, eight ordinal benefit/risk ratings and one binary choice. Both LLM arms and fitted controls receive the same 229 permitted Wave-1–3 demographic/personality history columns; target and semantically overlapping blocks are excluded. The pinned public source was prepared into fit/dev/cal/test respondents (1,000/200/200/658). The registered evaluation uses nominal/binary log loss and ordinal normalized RPS, with a pre-frozen equal-family normalized proper-risk aggregate,

CAL-only mixture selection, the full-data statistical comparator, and paired respondent inference. The transferred T0 is an instruction component, not the complete IPIP NEO genotype. Public Wave-4 labels and unknown model pretraining exposure prevent an unqualified untouched-holdout claim. No external result is inferred from source preparation.

The catalog stores one example text per question identifier. Our initial adapter treated the pricing identifiers as fixed products; this was an error in our use of the export, not a finding about the original study's forecasts. An outcome-blind audit of the source question objects found that all 2,058 respondents saw each of the 40 products once, but in randomized slots and at respondent-specific prices: 82,320 offers in total, including 7,490 zero-price offers. Each product appeared at 11 price levels; only 220 offers exactly matched the catalog example. The nine non-pricing question texts and ordered options were invariant. We therefore bind each forecast to the actual question text and the product's canonical identity rather than a QID9 presentation slot. A check restricted to the training split matched 43,000 of 43,000 stored question texts and 49,000 of 49,000 numeric slot choices to the corresponding source answers. The resulting R2 amendment was frozen before R2 fitting or GPU forecasts (SHA256 `1a968951...`); it supplies the same individual offer to both model arms and to the statistical controls. This repair followed an earlier, marginal-over-price calibration readback. R2 was specified before its own fitting and forecasts, after that earlier CAL analysis. The registered R2 comparison used 658 TEST respondents after the mixture weights were frozen on CAL. The crosswalk ledger and product-by-slot figure are in Supplement S3 (Section 3, Figure 1).

Fixed-weight readback. After E3 outcomes were inspected, we froze a CPU-only analysis of its *exposed* 600-person panel. For each representation we reuse the same ordinary ridge and set persona weight to 0.5 for both optimized and neutral forecasts, retaining each arm's historical calibration. Two paired neutral-minus-optimized RPS effects form a separate Holm family. We average the nine fixed per-person losses before respondent-level inference. This compares calibrated forecast packages at equal weight, not prompt text alone or new respondents. A fixed five-weight grid, three old-seed and leave-one-old-seed-out summaries, and uncalibrated forecasts are descriptive diagnostics (Supplement S4).

4.8 Additional diagnostic protocol

Saved-forecast diagnostics on the already exposed original cohort are post hoc. Separate Holm families cover 33 representation contrasts, nine objective contrasts and 48 contrasts per fitting-budget or mixture screen; they do not replace the original 110-test family. Linear CDF pooling follows $H = (1-w)R + wP$ Gneiting and Ranjan [2013]. Original Cal120 separates 60-person persona calibration from 60-person mixture selection; each rival selects its own weight. Training subsets overlap, and intervals condition on fixed statistical fits and old search seeds. Supplement S4 gives features, penalty/weight grids, all budgets, controls and full exposure history.

5 Results

5.1 Coverage and scoring verification

All 84 frozen final cells completed, each with 600 respondents and 36,000 observed target answers. Final receipts contain no failed requests or uncertain forecasts. They record 10,281,600 confirmed forecasts and 11,979,632,160 prompt tokens. Per cell, the count is $2(120)(60) + 3(600)(60) = 122{,}400$: calibration forecasts for base and evolved personas plus evaluation forecasts for base, evolved, and wrong personas. These totals concern the final stage, not search or total project cost.

A separate raw-prediction audit reloaded all 84 saved probability files and original item labels, selected calibration weights anew from saved calibration predictions, and reproduced each base, evolved, and wrong-persona per-respondent loss to absolute tolerance 10^{-10}. This validates the saved scoring computation; it does not independently replay model inference or the full search trajectories.

5.2 RQ1: Primary search and the full comparison

Table 2 reports the planned primary contrasts. The raw-fitness MAP-Elites-style method has no crossed interval wholly below zero against base on either model. Nor does it establish an advantage over OPRO on Qwen or native GigaEvo on Granite. These are inconclusive superiority comparisons, not evidence of equivalence.

Table 3 gives all 14 method–objective identities, including unfavorable and inconclusive outcomes. None improves over base under the stated rule on *both* models. Granite nevertheless has several base-beating methods: raw and calibrated GA, legacy-fitness GA, calibrated MAP-Elites-style search, ProTeGi, and best-of-budget. In particular, ProTeGi improves by 0.00250 and best-of-budget by 0.00171 RPS. A favorable base comparison therefore does not uniquely support evolutionary search or an archive mechanism. We do not select a new primary method from this table.

5.3 RQ2: Fitness-objective effects

The planned direct objective contrast for MAP-Elites-style search favors calibrated fitness on Granite: raw minus calibrated RPS is 0.00205 (crossed interval $[0.00015, 0.00340]$, Holm $p = .022$). Calibrated MAP-Elites-style search is also 0.00267 below base, and calibrated GA is 0.00266 below base. Both calibrated variants improve over their corresponding bases in all three observed Granite search

Table 2: Planned primary comparisons. Δ is raw-fitness MAP-Elites-style persona minus control in calibrated RPS (lower is better). Crossed 95% intervals resample respondents and search seeds; Holm-adjusted sign-flip p values condition on the three observed seeds. All rows use the same 600 people. Small discrepancies between rounded means and differences arise because differences are computed before rounding.

Model	Control	Persona RPS	Control RPS	ΔRPS [95% CI]	p_{H}
Qwen3.6	Base	0.13132	0.13008	$+0.00123$ $[-0.00053, +0.00363]$	.022
Qwen3.6	Wrong persona	0.13132	0.17410	-0.04278 $[-0.04781, -0.03696]$	.022
Qwen3.6	OPRO	0.13132	0.13152	-0.00020 $[-0.00143, +0.00083]$	1.000
Qwen3.6	Ridge	0.13132	0.10955	$+0.02177$ $[+0.01974, +0.02417]$	.022
Granite	Base	0.13594	0.13656	-0.00063 $[-0.00199, +0.00098]$	.023
Granite	Wrong persona	0.13594	0.16819	-0.03225 $[-0.03572, -0.02824]$	.022
Granite	Native GigaEvo	0.13594	0.13670	-0.00076 $[-0.00400, +0.00385]$	.022
Granite	Ridge	0.13594	0.10955	$+0.02639$ $[+0.02455, +0.02834]$	.022

Table 3: Complete method comparison: calibrated RPS difference from each cell's base, with crossed 95% intervals and respondent sign-flip Holm p values conditional on three fixed runs, from the same 110-test family. A dagger marks an interval wholly below zero and $p_{\mathrm{H}} < .05$. The first eleven rows use raw-RPS fitness; the final three are objective variants. Each entry aggregates three search seeds and 600 paired respondents.

	Qwen3.6		Granite	
Method	ΔRPS [95% CI]	p_{H}	ΔRPS [95% CI]	p_{H}
GA	$+0.00118$ $[+0.00053, +0.00178]$	.022	-0.00136 $[-0.00206, -0.00074]^{\dagger}$	.022
DE	$+0.00089$ $[-0.00007, +0.00201]$	.022	-0.00116 $[-0.00249, +0.00058]$	.022
PromptBreeder	$+0.00246$ $[-0.00069, +0.00495]$	.022	-0.00053 $[-0.00235, +0.00251]$	.023
MAP-Elites-style	$+0.00123$ $[-0.00053, +0.00363]$	.022	-0.00063 $[-0.00199, +0.00098]$	.023
GEPA-style	$+0.00010$ $[-0.00151, +0.00145]$	1.000	-0.00087 $[-0.00322, +0.00170]$	.022
Official GEPA	-0.00020 $[-0.00122, +0.00101]$	1.000	-0.00045 $[-0.00168, +0.00061]$	.842
Native GigaEvo	-0.00009 $[-0.00130, +0.00104]$	1.000	$+0.00014$ $[-0.00280, +0.00213]$	1.000
OPRO	$+0.00143$ $[-3.2 \times 10^{-6}, +0.00308]$	.022	$+0.00074$ $[-0.00136, +0.00420]$	.022
ProTeGi	$+0.00244$ $[+0.00101, +0.00405]$	.022	-0.00250 $[-0.00343, -0.00174]^{\dagger}$	.022
Random	$+0.00086$ $[+0.00039, +0.00133]$	.022	$+0.00115$ $[-0.00128, +0.00444]$	.022
Best-of-budget	-0.00044 $[-0.00112, +0.00020]$	.272	-0.00171 $[-0.00211, -0.00131]^{\dagger}$	.022
GA, calibrated fitness	-0.00044 $[-0.00104, +0.00016]$	.582	-0.00266 $[-0.00318, -0.00213]^{\dagger}$	.022
GA, legacy fitness	$+0.00080$ $[+3.0 \times 10^{-7}, +0.00164]$	.022	-0.00166 $[-0.00255, -0.00082]^{\dagger}$	.022
MAP-Elites, calibrated fitness	-0.00034 $[-0.00099, +0.00030]$	1.000	-0.00267 $[-0.00377, -0.00179]^{\dagger}$	.022

seeds (Supplement S4). Their Qwen intervals versus base cross zero. This differs from the direct objective effect: an objective can improve an otherwise harmful search without establishing a benefit over the neutral profile.

The additional direct GA analysis finds raw-minus-calibrated effects of 0.00163 on Qwen (crossed interval $[0.00111, 0.00219]$) and 0.00130 on Granite ($[0.00060, 0.00204]$). Each is positive in all three observed seeds and has conditional Holm $p = .002$ in the new nine-test family. The archival Part B objective plot shows all four contrasts, including the Qwen MAP-Elites seed with an unfavorable effect. All five interaction intervals cross zero: significance in one subgroup and not another does not demonstrate an interaction. The same MAP-Elites Granite contrast has $p = .022$ in the original 110-test family and $p = .002$ in the additional nine-test family; the smaller value reflects a family change, not new evidence.

Candidate rankings reveal why raw RPS is not interchangeable with the calibrated objective. Across the 24 search pools, median Spearman correlation is 0.344 (range -0.562 to 0.788), median pairwise discordance is 35.6% (20.7–67.4%), and median top-five overlap is zero (range 0–4). Thus calibration changes candidate ordering, not only a score's scale. This is a descriptive search-panel result; it does not prove that ranking changes caused held-out gains. Final raw-RPS shortlist selection remains unchanged.

5.4 Additional-respondent transfer of frozen templates

All 24 fixed-template cells completed on the additional 600 respondents; source, attempt, panel, probability, cost, and serving-log checks passed. An independent CPU readback reproduced all saved losses and 16 reported contrasts from raw forecasts (maximum absolute difference 8.33×10^{-17}). In the declared Granite three-test family, calibrated MAP-Elites-style and GA templates beat their respective neutral bases by 0.00216 RPS (crossed 95% interval $[0.00134, 0.00311]$) and 0.00229 ($[0.00166, 0.00289]$); both have Holm-adjusted sign-flip $p = .0006$. The calibrated-

versus-raw MAP-Elites-style effect is +0.00152 raw minus calibrated, but its crossed interval $[-0.00018, 0.00273]$ includes zero, so it *does not* pass the combined rule despite adjusted $p = .0006$ and two favorable seed effects. This is partial fixed-template transfer, not a replication of that objective effect under the declared rule (full table in Supplement S4).

Qwen calibrated GA and MAP-Elites-style winners beat their raw-fitness counterparts by 0.00181 and 0.00168 RPS in diagnostic contrasts, but neither beats its neutral base under the interval rule; both intervals include zero (full diagnostic estimates in Supplement S4). The same-history ridge scores 0.10821 on these people, against 0.12906/0.12920 for calibrated Qwen GA/MAP and 0.13340/0.13353 for calibrated Granite GA/MAP. Thus the positive Granite base comparisons do not establish practical superiority to the stronger statistical control. Historical winners, old calibration and three old search seeds were fixed; the test neither evaluates new search trajectories nor the developmental aligned-selector winners.

5.5 Exploratory selector intervention

Fresh same-shortlist scoring changed 11 of 12 GA winners. Effects below are calibrated minus raw selector RPS; negative favors calibration. For Qwen calibrated search the effect is -0.000153, with crossed 95% interval $[-0.00050, 0.00018]$; for raw search it is -0.000587, with interval $[-0.00146, 0.00004]$. For Granite calibrated search the effect is -0.000013, with interval $[-0.00021, 0.00030]$; for raw search it is -0.000473, with interval $[-0.00133, 0.00010]$. All four intervals include zero, despite adjusted sign-flip $p = .0020$ for the two raw-search effects. Thus substantial winner turnover does not establish final-score benefit under our combined rule. The sole search-objective effect with an interval below zero is Qwen calibrated versus raw search under the *raw* selector; it does not show a calibrated-selector gain. The archival Part B ten-test table gives all seed effects and adjusted probabilities.

5.6 Frozen Granite-80 mixture transfer

The third panel's four declared comparisons all pass the fixed rule (Table 4). The optimized mixture scores 0.13236 RPS with categorical features and 0.13153 with raw history. The corresponding neutral mixtures score 0.13382/0.13420 and 0.13278/0.13328, respectively. These are modest gains for this old Granite recipe, conditional on nine fixed combinations; they do not identify the isolated effect of changing a prompt. The full 8,000-person same-history ridge scores 0.10995 and remains substantially better than either hybrid. All 216,000 new forecasts were confirmed; an independent CPU recomputation agreed with saved losses and statistics within 9.38×10^{-11}, below the fixed 10^{-10} tolerance. Individual combination results, components, weights, exposure limits and costs appear in Supplement S4.

Table 4: Third-panel Granite-80 old-recipe transfer. Positive gain is comparator minus optimized hybrid RPS; paired intervals condition on nine fixed old combinations. All four Holm-adjusted $p = .0008$. N=600.

Features	Neutral comparator	Gain [paired 95% CI]
Categorical	Prior-anchored	0.001465 [0.001213, 0.001722]
Categorical	Same ordinary ridge	0.001843 [0.001595, 0.002088]
Raw	Prior-anchored	0.001251 [0.001016, 0.001490]
Raw	Same ordinary ridge	0.001747 [0.001519, 0.001974]

5.7 Same-panel B1: fixed hybrid versus selected baselines

Figure 1 reports every frozen fitting budget. The two prespecified n=8,320 primary differences are *against* the hybrid: +0.021679 RPS for raw history (conditional paired 95% CI [0.020556, 0.022857]) and +0.018311 for the categorical profile ([0.017188, 0.019426]). Holm-adjusted $p = 0.0004$ for each, but both practical-benefit gates are false because their direction is positive. At n=80 the fixed hybrid scores better in both representations under its larger historical supervision. At n=160 raw is inconclusive (CI crosses zero) while categorical favors the hybrid; from n=320 every examined point favors the selected baseline. These are secondary, exposed-panel contrasts, not an estimated first crossing or a literal questionnaire answer-equivalent. The earlier full-8,000 ridge RPS 0.10995 versus raw hybrid 0.13153 already showed the broad gap; the new same-panel, within-budget CV family comparison tests that conclusion with capacity selection and both representations.

The accompanying aggregate CSV reports all twelve signed differences, intervals and adjusted tests. The same rows and their conditional estimand appear in Supplement S3, Table 1.

How many observed answers match the fixed hybrid? We held the 8,320 training respondents and the 60 target items fixed. For each of three frozen orders of the 60 input items we fitted statistical predictors to nested prefixes of 0, 1, 2, 5, 10, 20, 30, 40 and 60 answers. Three-fold training-respondent cross-validation selected the predictor at each prefix. We scored all predictors on the same exposed E3 panel of 600 people. The reference RPS, 0.131532, averages losses across the nine fixed historical raw-hybrid fits, each using all 60 input answers.

The decoder's mean RPS across orders decreased from 0.151894 with zero answers to 0.132718 with ten, 0.124497 with twenty and 0.109854 with all sixty. Every order yielded a monotone raw curve. Linear interpolation between ten and twenty answers gives $m^* = 11.44$ for the average curve. All 2,000 paired respondent-bootstrap replicates crossed within the grid (95% interval [10.46, 12.38], conditional on the fixed fits and item orders). Individual orders crossed at 9.17, 10.99 and 14.19 answers. Jointly resampling respondents and the three observed orders gives a descriptive interval of [9.13, 13.88]. The crossing mea-

sures equivalent predictive loss for this decoder and target. Prospective item selection and validation in new people would be needed to justify a shorter questionnaire.

Equal-weight diagnostic on the exposed E3 panel. At weight 0.5 with the same ordinary ridge, the old calibrated optimized forecast package beats its neutral counterpart by 0.001477 RPS for categorical features (paired 95% CI $[0.001291, 0.001676]$) and 0.001448 for raw history ($[0.001264, 0.001643]$); both Holm-adjusted $p = .0004$. All three historical-seed and all leave-one-seed-out mean differences are positive. This post hoc readback removes the E3 mixture-weight difference but preserves arm-specific calibration and reuses the same exposed people; it is neither a prompt-only effect nor fresh confirmation. The uncalibrated half-weight gains are larger but both absolute scores are worse. Full grid, controls, and independent raw-loss verification appear in Supplement S4.

5.8 External instrument: held-out calibration of the statistical control

In an unregistered auxiliary analysis, we evaluated the frozen R2 statistical control on the 200 calibration respondents excluded from its 1,200-person fit. Each respondent supplied one observed answer for each of the 49 endpoints. For each canonical product, selection within the training data considered the marginal prior, the posted price, the permitted history and the two sources combined. All 40 selected models used the actual posted price (21 price-only and 19 price-plus-history; 18 logistic and 22 tree models). The comparator is the product-specific, Laplace-smoothed response prior fitted on those same 1,200 people.

Across the 40 hypothetical purchase choices, mean log loss fell from 0.637315 (prior) to 0.519243 (selected decoder). The paired difference was -0.118072 (descriptive 95% respondent-bootstrap interval $[-0.133371, -0.102856]$); 39 products improved and one worsened. This contrast measures the fitted predictor's performance against a product-specific prior. Isolating the contribution of personal history requires a direct comparison with a price-only forecast.

Across the eight ordinal ratings, mean RPS was 0.121872 for the decoder and 0.121223 for the prior (difference $+0.000649$, 95% interval $[-0.001265, 0.002567]$); four ratings improved, three worsened and one tied. The marble-choice decoder selected the prior, with log loss 0.643622 for both. Intervals resample 200 people, conditional on the fitted models, averaging endpoints of the same type within person. The 1,400-person fit includes CAL and is reserved for the registered TEST comparison.

In the earlier, slot-indexed analysis, a history-only decoder selected the marginal forecast on 45 of 49 endpoints. At the 40 price slots, decoder and prior both scored 0.684166. Each slot pooled different products and prices across people. The R2 product-indexed scores use different target definitions, so a numerical difference between these two analyses cannot estimate an effect of the repair.

Feasibility of stimulus-matched model forecasts. On ten development respondents, Granite and Qwen each produced all 980 probability vectors (49 endpoints, two instructions). Across the 980 calls per model, the minimum total probability mass on valid answer tokens before renormalization was 0.9954 for Granite and 0.9716 for Qwen, above the frozen 0.5 threshold. Each model used one H200, for 56 min 24 s and 45 min 1 s respectively, processing 24.98 and 24.86 million input tokens. These checks use no outcome labels. They admitted both models to the 200-person calibration stage.

Held-out prediction on the external survey. All 658 TEST respondents had labels for all 49 endpoints. Relative to the statistical decoder fitted on 1,400 respondents, standalone LLMs and fixed half-weight mixtures had higher normalized loss in all eight comparisons. Both CAL-selected Granite mixtures improved modestly: $\Delta = -0.00404$ for neutral (95% CI $[-0.00723, -0.00096]$, Holm $p = .0360$) and -0.00384 for T0 ($[-0.00692, -0.00085]$, Holm $p = .0486$). Both CAL-selected Qwen mixture intervals included zero. In a post hoc direct paired comparison of the two Granite mixtures, the same CAL rule gave T0 minus neutral $\Delta = +0.000197$ (95% CI $[-0.000541, 0.000957]$). Both arms used the same calibration respondents, weight grid, loss and tie rule, with weights selected separately. At a common fixed weight of 0.5, the contrast was -0.014990 (95% CI $[-0.017155, -0.012732]$). The fixed-weight comparison favors T0. The CAL-selected contrast was near zero and its paired interval included zero. These exploratory contrasts reuse R2 TEST and are reported separately from the registered Holm families. In the standalone comparisons, T0 versus neutral reduced Granite loss by 0.05703 but increased Qwen loss by 0.03066. These are dimensionless differences normalized by endpoint-specific CAL prior loss, with equal weight for the two endpoint families. Intervals condition on the frozen forecasts. The complete results are in Supplement S3. In a separate exploratory decomposition on the same R2 TEST, the frozen CAL-selected Granite mixtures also improved on their own decoder1200 component: neutral -0.002886 (95% interval $[-0.004478, -0.001282]$), T0 -0.002689 ($[-0.004239, -0.001123]$). Both Qwen intervals included zero. Decoder1200 minus Full1400 was -0.001150 ($[-0.004417, 0.002007]$). Intervals are descriptive, conditional on fixed forecasts and outside the registered Holm families. The comparison holds the fitted decoder1200 fixed. The mixtures additionally use 200 CAL respondents for weight selection and require LLM forecasts. Full1400 fits on those CAL respondents as well as the original 1,200. The separate neutral contrast tests the two instruction-specific mixtures.

In the separate exploratory comparison on the 40 pric-

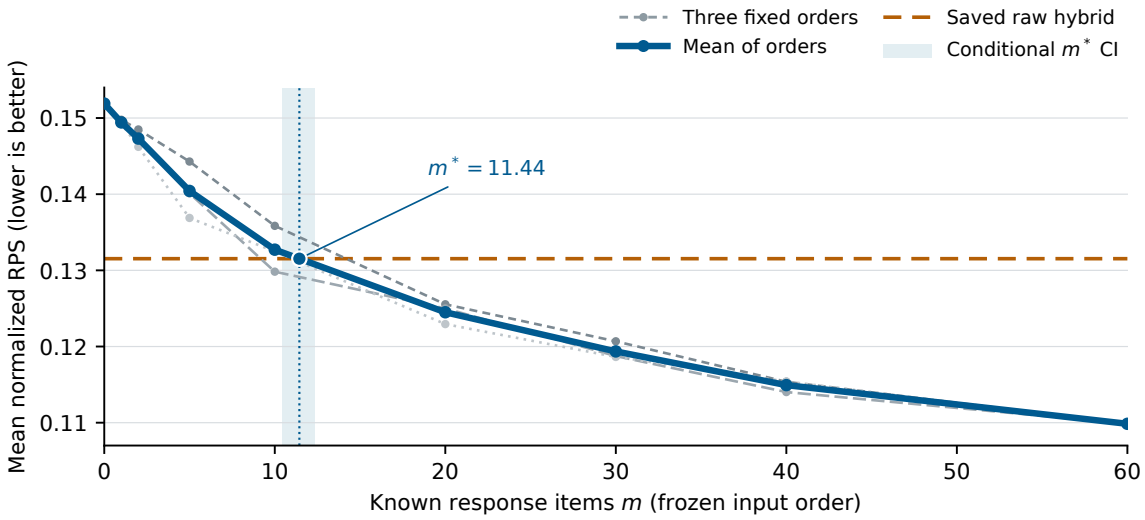


Figure 3: Observed answers needed to match the fixed hybrid on the exposed E3 panel. Thin lines show three frozen orders of input items. The thick line averages the decoder's losses across orders at each measured budget. The horizontal line marks the mean RPS of the nine historical raw-hybrid fits. The crossing interval resamples respondents, conditional on the three orders and fitted models.

ing products, adding the respondent's history to price increased mean log loss by 0.00420 (95% paired respondent interval $[-0.00015, 0.00852]$). The interval includes zero. This comparison uses the same 658 people and fixed statistical fits.

5.9 RQ3: Profile matching and prediction

The primary template instantiated with the matched profile outperforms the same template instantiated with a deranged respondent profile by 0.04278 RPS on Qwen and 0.03225 on Granite (Table 2). This comparison detects predictive signal from matching the supplied profile to its respondent; it does not test unique identification. It also does not isolate optimization, because neutral templates receive matched histories too.

The selection-chosen ridge predictor has held-out RPS 0.10955, compared with 0.13132 and 0.13594 for the primary Qwen and Granite personas. Ridge also outperforms both calibrated-fitness variants. For descriptive context, the prior scores 0.15174 and ExtraTrees 0.13127; evaluation scores are not used to change the selected decoder. The ridge comparison concerns the full prediction pipelines, given their different representations and fitting procedures.

Ridge using categorical profiles scores 0.11384; raw-history ridge scores 0.10955. Categorical-profile ridge retains 89.8% of raw-history ridge's prior-relative prediction gain. Every LLM arm is worse than categorical ridge in the 33-test family (all 28 relevant intervals positive, Holm $p < .01$). This is an observed prediction-gain ratio, not a fraction of human information or a causal compression decomposition. Ridge uses 8,000 labelled person-level mappings; the LLM uses those labels for priors plus search, selection and calibration targets. Supervision and learning procedures are not matched. Full representation contrasts and base-to-ridge descriptive ratios are in Supplement S4.

5.10 Exploratory budget and mixture diagnostics

On the exposed original cohort, calibrated GA can beat shared-prior ridge at 80 fitting people, while ridge has lower mean RPS at 320. Against raw-history ridge on exact Search80, Qwen gains 0.00530 RPS (crossed interval $[0.00422, 0.00638]$, conditional Holm $p = .0096$); Granite does not pass the combined rule. These fitting counts omit selection, calibration and prior supervision. Archival Part B learning curves and answer-equivalent diagnostics retain all budgets and their conditional uncertainty.

Mixtures beat both components in all 12 original conditions, but neutral controls restrict optimized-recipe utility. Against neutral plus anchored ridge, only Granite categorical/raw 80 pass the practical/statistical rule; Qwen loses in all six conditions, five significantly. This adaptive 24-test screen motivated E3, not a general optimizer claim. Full contrasts, weights, component gains and the separate 48-test family remain in Supplement S4.

6 Discussion

E1 supports historical calibrated Granite templates over neutral, but its MAP objective interval crosses zero; Qwen base advantages remain unsupported. Calibration changes

ranks and most same-shortlist winners without a clear selector benefit. Retaining historical winners precludes end-to-end optimizer replication or uniform transfer claims.

The third-panel Granite-80 recipe beats both neutral controls in categorical and raw representations. The ordinary-ridge control fixes coefficients but changes the old-selected weight with the prompt, supporting the recipe rather than a prompt-only or new-search effect. E4's smaller positive equal-weight package difference rules out unequal weights alone; it remains post hoc with distinct calibration.

The B1 comparison changes the practical interpretation. With the same 60 observed test-person items and a separately tuned statistical capacity at each fit budget, both baseline representations overtake the fixed hybrid by n=320 and are much better at n=8,320. The small-budget advantage of the hybrid survives as an observed point, with broader historical supervision than the baseline fit budget counts. These data do not identify the causal value of prompt text, nor a literal reduction in questions asked. A categorical-profile baseline also wins at large n, so uncompressed raw inputs alone cannot explain the full-data gap. Which component of the old mixture helps at small n remains a narrower question than whether a persona adds value over the strongest available same-history predictor.

7 Limitations and reproducibility

A separate calibration-only replay with pinned inputs, sources and serving settings failed a fixed numerical compatibility rule on one of two Qwen repeats; both Granite repeats passed (Supplement S4). This is not an optimizer test or a diagnosis of the cause of variation. Two subsequent full selected-template replays retain the GA objective contrast on this exposed cohort: raw-minus-calibrated RPS is 0.00161/0.00159 for Qwen and 0.00130/0.00130 for Granite. All four crossed intervals exclude zero (within-family Holm $p = 0.0008$; Supplement S4). This supports the contrast's robustness in these executions, but does not repair pointwise compatibility or establish generalization over serving environments. The selector study's strict selection audit failed on a post-receipt server teardown diagnostic; a versioned gate accepted that diagnostic while verifying the saved selection arrays. The final gate and independent CPU readback passed. Fresh Qwen neutral-base forecasts nevertheless have per-cell mean absolute drift of about 0.00031–0.00037 RPS from the archived base (maximum individual drift 0.0028–0.0048); Granite drift is much smaller. These checks do not repair pointwise serving compatibility. The planned E5 replication with new optimizer seeds remains unexecuted: its Granite readout failed the unchanged 5×10^{-4} pointwise gate against E5's frozen calibrated-GA reference.

Evaluation covers one inventory, two readout families and existing-corpus participants; public-data pretraining exposure is unknown. E1 and E3 are disjoint from each other and the reachable inventoried exposure union, but share a corpus and item family; global lifetime exposure and untouched measurement choices remain unproven. Their targets did not tune the frozen recipes. Transfer fixes old prompts and calibration, so it does not test new optimizer runs, alternate selector winners or external behavior.

E3 fixes three old search seeds, three ridge subsets and old-selected weights; its intervals condition on those choices. Its protocol followed earlier-outcome inspection. Neither fit-$n = 80$ nor a new respondent panel implies equal total supervision or external population validation. E4 reuses E3's exposed labels and forecasts, adding neither a second transfer panel nor independent old-seed replication. Other post hoc specifications preceded execution but their questions followed original outcomes. Three seeds limit inference over optimizer randomness; pointwise bootstrap intervals are not simultaneous across methods.

B1 reuses the now-exposed E3 panel and averages nine correlated old hybrid combinations against three fixed baseline subsets. Its respondent intervals are conditional on these runs; Holm-adjusted sign-flip tests do not turn this into a fresh confirmation. The n=8,320 contrast compares total unique task-labelled people touched somewhere in the historical LLM recipe to a statistical fitting budget, not identical supervision semantics or compute. The Twin-2K-500 contrasts condition on the two fixed model revisions, instructions, fitted decoders and CAL weights. R2 was frozen after inspection of the earlier history-only CAL diagnostic but before its own fitting and GPU forecasts. Its exact-stimulus repair resolves a measurement mismatch, not the risk of model pretraining exposure or post-CAL design influence. Wave-4 answers are public, and hypothetical choices/ratings do not validate real purchases or behavior.

Equal forecast budgets do not imply equal total cost. The recorded historical historical minimum is 300.60 allocated GPU-hours. Failed/cancelled request work, other development, mutator/completion tokens, CPU work and dollars remain unresolved. The separate replays omit search, prior fitting and old calibration costs; forecast/token/allocation accounting appears in Supplement S4 (Tables S15–S16, combined PDF p. 18). No end-to-end efficiency claim follows. Native GigaEvo is private: reproduction requires access or a separately validated substitute, which must not be labelled native reproduction.

Manifests, model locks, templates, probabilities, calibration and loss receipts are retained. Human histories and exposure IDs require de-identification and source permissions; labelled training humans count toward cost. Shared data/methodology do not establish an independent publication contribution.

AI assistance. This revision additionally used Claude Code for independent manuscript critique.

Codex assisted inspection, retrieval, design, analysis and writing; Claude supplied generic plotting/layout designs. Scientific labels and measurements were bound

locally; figures were checked against recorded outputs. Supplement S4 records available versions, prompts and historical-provenance limits.

8 Conclusion

The frozen Granite calibrated-GA persona-plus-ridge recipe beats two neutral mixtures on a third, same-corpus panel. Its practical comparison with a capacity-selected same-history predictor is less favorable: at 8,320 fitted people the hybrid has 0.021679 higher raw-history RPS and 0.018311 higher profile-history RPS, with both declared practical gates false. Both baseline representations lead from 320 fitted people onward; the low-budget hybrid advantage is conditional on its old, larger supervision pipeline. With training size fixed at 8,320, a raw-history decoder reaches the 60-answer hybrid's mean loss at an interpolated 11.44 known answers on the same panel. The original planned optimizer comparison is null; E1's Granite neutral gains and E2/E4 diagnostics retain their narrower interpretations. We therefore measure a boundary on persona value, rather than claim general superiority, sample-efficiency savings, or behavioral transfer. The external-protocol validity audit recovered each respondent's randomized product and price and bound both model arms and controls to that same stimulus. In that survey, adaptive mixing gave a small Granite gain over the full-data control for both instructions. The opposing standalone T0 effects across Granite and Qwen limit any claim of portable persona benefit. The direct CAL-selected Granite comparison yielded a near-zero estimate (Section 5.8).

References

Lakshya A. Agrawal, Shangyin Tan, Dilara Soylu, Noah Ziems, Rishi Khare, Krista Opsahl-Ong, Arnav Singhvi, Herumb Shandilya, Michael J. Ryan, Meng Jiang, Christopher Potts, Koushik Sen, Alexandros G. Dimakis, Ion Stoica, Dan Klein, Matei Zaharia, and Omar Khattab. GEPA: Reflective prompt evolution can outperform reinforcement learning. In *International Conference on Learning Representations*, 2026. URL `https://arxiv.org/abs/2507.19457`.

Gati V. Aher, Rosa I. Arriaga, and Adam Tauman Kalai. Using large language models to simulate multiple humans and replicate human subject studies. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 337–371, 2023. URL `https://proceedings.mlr.press/v202/aher23a.html`.

Lisa P. Argyle, Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. Out of one, many: Using language models to simulate human samples. *Political Analysis*, 31(3):337–351, 2023. doi: 10.1017/pan.2023.2.

Farima Fatahi Bayat, Moin Aminnaseri, Pouya Pezeshkpour, and Estevam Hruschka. Reflective prompt tuning through language model function-calling, 2026. URL `https://arxiv.org/abs/2605.21781`.

Marcel Binz, Elif Akata, Matthias Bethge, Franziska Brändle, Fred Callaway, Julian Coda-Forno, Peter Dayan, Can Demircan, Maria K. Eckstein, Noémi Éltető, Thomas L. Griffiths, Susanne Haridi, Akshay K. Jagadish, Li Ji-An, Alexander Kipnis, Sreejan Kumar, Tobias Ludwig, Marvin Mathony, Marcelo Mattar, Alireza Modirshanechi, Surabhi S. Nath, Joshua C. Peterson, Milena Rmus, Evan M. Russek, Tankred Saanum, Johannes A. Schubert, Luca M. Schulze Buschoff, Nishad Singhi, Xin Sui, Mirko Thalmann, Fabian J. Theis, Vuong Truong, Vishaal Udandarao, Konstantinos Voudouris, Robert Wilson, Kristin Witte, Shuchen Wu, Dirk U. Wulff, Huadong Xiong, and Eric Schulz. A foundation model to predict and capture human cognition. *Nature*, 644:1002–1009, 2025. doi: 10.1038/s41586-025-09215-4.

James Bisbee, Joshua D. Clinton, Cassy Dorff, Brenton Kenkel, and Jennifer M. Larson. Synthetic replacements for human survey data? the perils of large language models. *Political Analysis*, 32(4):401–416, 2024. doi: 10.1017/pan.2024.5.

Chrisantha Fernando, Dylan Sunil Banarse, Henryk Michalewski, Simon Osindero, and Tim Rocktäschel. Promptbreeder: Self-referential self-improvement via prompt evolution. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 13481–13544. PMLR, 2024. URL `https://proceedings.mlr.press/v235/fernando24a.html`.

Leon Fröhling, Jens Rupprecht, Markus Strohmaier, and Claudia Wagner. When persona attributes improve population alignment in large language models, 2026. URL `https://arxiv.org/abs/2609.02526`.

Tilmann Gneiting and Adrian E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477):359–378, 2007. doi: 10.1198/016214506000001437.

Tilmann Gneiting and Roopesh Ranjan. Combining predictive distributions. *Electronic Journal of Statistics*, 7:1747–1782, 2013. doi: 10.1214/13-EJS823. URL `https://doi.org/10.1214/13-EJS823`.

Qingyan Guo, Rui Wang, Junliang Guo, Bei Li, Kaitao Song, Xu Tan, Guoqing Liu, Jiang Bian, and Yujiu Yang. Connecting large language models with evolutionary algorithms yields powerful prompt optimizers. In *International Conference on Learning*

Representations, 2024. URL https://proceedings.iclr.cc/paper_files/paper/2024/hash/9156b0f6dfa9bbd18c79cc459ef5d61c-Abstract-Conference.html.

Tiancheng Hu and Nigel Collier. Quantifying the persona effect in LLM simulations. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10289–10307, Bangkok, Thailand, 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.554. URL https://aclanthology.org/2024.acl-long.554/.

Nusrath Jahan and Johnathan Mell. Unraveling the tapestry of deception and personality: A deep dive into multi-issue human-agent negotiation dynamics. In *Proceedings of the 23rd International Conference on Autonomous Agents and Multiagent Systems*, pages 916–925. IFAAMAS, 2024. doi: 10.65109/sjpk9444. URL https://www.ifaamas.org/Proceedings/aamas2024/pdfs/p916.pdf.

Mumin Jia, Yilin Chen, Divya Sharma, and Jairo Diaz-Rodriguez. When can digital personas reliably approximate human survey findings?, 2026. URL https://arxiv.org/abs/2605.10659v3. Accepted at NeurIPS 2026; arXiv version 3, 13 October 2026.

John A. Johnson. Measuring thirty facets of the five factor model with a 120-item public domain inventory: Development of the IPIP-NEO-120. *Journal of Research in Personality*, 51:78–89, 2014. doi: 10.1016/j.jrp.2014.05.003. URL https://doi.org/10.1016/j.jrp.2014.05.003.

Jean-Baptiste Mouret and Jeff Clune. Illuminating search spaces by mapping elites, 2015. URL https://arxiv.org/abs/1504.04909.

Lewis Newsham and Daniel Prince. Personality-driven decision-making in LLM-based autonomous agents. In *Proceedings of the 24th International Conference on Autonomous Agents and Multiagent Systems*, pages 1538–1547. IFAAMAS, 2025. URL https://ifaamas.org/Proceedings/aamas2025/pdfs/p1538.pdf.

Joon Sung Park, Carolyn Q. Zou, Jonne Kamphorst, Niles Egan, Aaron Shaw, Benjamin Mako Hill, Carrie Cai, Meredith Ringel Morris, Percy Liang, Robb Willer, and Michael S. Bernstein. LLM agents grounded in self-reports enable general-purpose simulation of individuals, 2026. URL https://arxiv.org/abs/2411.10109v3. Version 3, 28 June 2026; first posted in 2024.

Reid Pryzant, Dan Iter, Jerry Li, Yin Tat Lee, Chenguang Zhu, and Michael Zeng. Automatic prompt optimization with "gradient descent" and beam search. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 7957–7968. Association for Computational Linguistics, 2023. doi: 10.18653/v1/2023.emnlp-main.494. URL https://aclanthology.org/2023.emnlp-main.494/.

Mikayel Samvelyan, Davide Paglieri, Minqi Jiang, Jack Parker-Holder, and Tim Rocktäschel. Multi-agent diagnostics for robustness via illuminated diversity. In *Proceedings of the 23rd International Conference on Autonomous Agents and Multiagent Systems*, pages 1630–1644. IFAAMAS, 2024. doi: 10.65109/OPQC5611. URL https://www.ifaamas.org/Proceedings/aamas2024/pdfs/p1630.pdf.

Lindia Tjuatja, Valerie Chen, Tongshuang Wu, Ameet Talwalkwar, and Graham Neubig. Do LLMs exhibit human-like response biases? a case study in survey design. *Transactions of the Association for Computational Linguistics*, 12:1011–1026, 2024. doi: 10.1162/tacl_a_00685. URL https://aclanthology.org/2024.tacl-1.56/.

Olivier Toubia, George Z. Gui, Tianyi Peng, Daniel J. Merlau, Ang Li, and Haozhe Chen. Database report: Twin-2k-500: A data set for building digital twins of over 2,000 people based on their answers to over 500 questions. *Marketing Science*, 44(6):1446–1455, 2025. doi: 10.1287/mksc.2025.0262. URL https://doi.org/10.1287/mksc.2025.0262.

Weilai Xu, Fred Charles, and Charlie Hargood. Generating stylistic and personalized dialogues for virtual agents in narratives. In *Proceedings of the 22nd International Conference on Autonomous Agents and Multiagent Systems*, pages 737–746. IFAAMAS, 2023. URL https://www.ifaamas.org/Proceedings/aamas2023/pdfs/p737.pdf.

Chengrun Yang, Xuezhi Wang, Yifeng Lu, Hanxiao Liu, Quoc V. Le, Denny Zhou, and Xinyun Chen. Large language models as optimizers. In *International Conference on Learning Representations*, 2024. URL https://proceedings.iclr.cc/paper_files/paper/2024/hash/3339f19c5fcee3ad74502947a32be9e6-Abstract-Conference.html.