

A Constant Vector Beats Your Digital Twin: On the Evaluation of Cluster-Level LLM Personality Simulation

Anonymous ACL submission

Abstract

The standard way to evaluate an LLM “digital twin” of a human personality group — agreement between the model’s questionnaire answers and real respondents’ — rewards the wrong thing. For this answer-similarity metric we prove that a simulator reproducing the target population’s response distribution *exactly* can never beat a constant vector, and is strictly beaten unless every item’s mass sits on its medians; the gap is the difference between the Gini mean difference and the mean absolute deviation about the median. On 410,168 IPIP-NEO-120 respondents the inequality is strict on all 480 cluster–item combinations and the penalty is 7.75 points. Empirically, across five LLMs and four personality clusters, a per-item cluster mean scores 0.796 while evolved personas average 0.739, losing in 20 of 20 cells; a *real human drawn from the target cluster* scores 0.721 and also loses to the constant in 20 of 20 cells. The released benchmark also scores raw model self-reports against reverse-recorded human answers on 55 of 120 items; we correct this throughout. A metric that ranks a constant above a genuine group member cannot be measuring personality simulation. We propose a replacement suite — a classifier two-sample test, per-item distributional fidelity, and a variance ratio — each reported against an explicit human ceiling and constant floor, and show it inverts the ranking correctly. Under it, current personas are near-perfectly separable from humans (C2ST-AUC 0.994 vs. a human ceiling of 0.564) and reach only 0.43 of human response *spread* as a standard-deviation ratio. Correct persona assignment improves CRPS for all eleven models, yet all lose to a cluster empirical prior fitted on other respondents, and an exact decomposition shows the excess loss of the model’s mixture dominates the shortfall. We release code and per-cell results.

1 Introduction

LLM simulations of human respondents are validated by administering a psychometric inventory and scoring the agreement between the model’s answer vector and a human’s. That scoring step invalidates the conclusion it is used to draw. Our starting point is a concrete artefact. In a cluster-level digital-twin system built on IPIP-NEO-120 (Johnson, 2014), evolutionary prompt optimisation raises answer-level agreement while simultaneously *degrading* the recovered Big Five profile — mean absolute error, trait similarity and facet similarity all move the wrong way. That pattern is usually reported as an open puzzle. We show the metric has a degenerate optimum, and measure where that bites.

Our contributions are:

1. **A proposition (§3):** under answer-similarity, a simulator that faithfully samples the target population’s response distribution is *dominated* by the best constant predictor — strictly so on every IPIP-NEO item we measure. The penalty is the excess of the Gini mean difference over the mean absolute deviation about the median, divided by the scale range.
2. **An empirical audit (§4):** across five LLMs $\times$ four clusters, on the original system’s own train/test splits, the matched cluster-mean baseline beats every model in every cell, and a global mean that ignores the clusters entirely beats it in 19 of 20.
3. **A diagnosis of the human ceiling (§7):** a real respondent from the target cluster — the ideal a cluster-level twin can aspire to — also loses to the constant in every cell.
4. **A replacement suite (§8):** C2ST-AUC, distributional fidelity and variance ratio, each anchored by a human ceiling and a constant

floor. It ranks humans above constants, as any valid measure must.

5. **A cluster-validity audit** (§12): the $k=4$ partition these systems condition on is not supported by any internal criterion.

2 Setting

Task. A population of respondents is partitioned into K personality clusters by k -means over the 30 IPIP-NEO facet scores. For each cluster a natural-language description is used to condition an LLM, which then answers the 120 IPIP-NEO items on a 1–5 Likert scale. Prompt wording is optimised by an evolutionary search whose fitness is answer agreement with held-out human members of the cluster.

Two regimes appear here and are not interchangeable. The *cluster* regime conditions on the centroid alone, giving every member one prompt; it is what Prop. 1 and the 7.75-point penalty address. The *individual* regime, used by the released prompt builder and by §11, inserts each respondent’s own scores as intensity modifiers. By Remark 1 one constant is optimal in the first and a per-respondent conditional median in the second; we name the regime wherever a result depends on it.

The metric. A simulated answer vector x is scored against a human answer vector y by

$$S(x, y) = \frac{1}{J} \sum_{j=1}^J \left(1 - \frac{|x_j - y_j|}{4}\right), \quad (1)$$

with $J=120$ and the divisor 4 the scale range. Cell scores average S over the held-out respondents. Equation (1), or a monotone transform of it, is the objective the search optimises and the number reported.

3 A constant is optimal under S

Fix a cluster and an item j , and let P_j be the cluster’s response distribution over $\{1, \dots, 5\}$. Write $Y \sim P_j$ for a randomly drawn human respondent.

Proposition 1. *Let m_j be a median of P_j , let $\text{MAD}_j = \mathbb{E}_Y |m_j - Y|$ and let $\text{GMD}_j = \mathbb{E}_{X,Y} |X - Y|$ for $X \sim P_j$ independent of Y (the Gini mean difference). Then*

- (a) *the best constant predictor attains expected score $1 - \text{MAD}_j/4$;*

- (b) *a faithful simulator, i.e. one whose output is distributed as P_j , attains expected score $1 - \text{GMD}_j/4$;*

- (c) *$\text{GMD}_j \geq \text{MAD}_j$, with equality iff X lies almost surely in the median set of P_j .*

Hence the constant is at least as good as the faithful simulator, and strictly better unless P_j places all its mass on medians.

Remark 1 (Where this applies). *Proposition 1 is unconditional; the audited systems condition on something. Let C be everything the simulator observes, its draw independent of the response given C . Under (1) an optimal simulator emits a conditional median $m_j(C)$, so one constant stays optimal exactly when the conditional median sets share a value almost surely, item by item. Cluster-only conditioning thus admits a cluster-specific constant optimum; individually informative conditioning can make different deterministic answers optimal, so the constant need not win. A deterministic conditional median is always optimal and a randomised predictor ties exactly when its support lies in that median set, so absolute error never uniquely rewards faithful conditional sampling; richer conditioning moves which prediction is rewarded.*

Proof. (a) and (b) are immediate from (1). For (c),

$$\begin{aligned} \text{GMD}_j &= \mathbb{E}_X [\mathbb{E}_Y |X - Y|] \\ &\geq \mathbb{E}_X [\min_c \mathbb{E}_Y |c - Y|] \\ &= \mathbb{E}_Y |m_j - Y| = \text{MAD}_j, \end{aligned}$$

since the median minimises expected absolute deviation. Equality requires X to attain that minimum almost surely, i.e. to lie in the median set of P_j . This is weaker than degeneracy: a balanced two-point law $P(Y=1) = P(Y=5) = \frac{1}{2}$ has $\text{GMD} = \text{MAD} = 2$, so the bound is tight there despite large variance. On the IPIP items it is strict throughout (Table 2). $\square$

The consequence is not a technicality. Under S the behaviour we want — reproducing the population’s response distribution — is penalised relative to one fixed vector, so the optimum is degenerate in exactly the variance that carries a profile. Whether an optimiser is driven there is empirical and not uniform: S_0 ships a near-worst model from a heterogeneous pool (App. D), yet prompt search under it improves proper scores too (App. E).

3.1 The metric is CRPS with its dispersion term deleted

Proposition 1 is not an accident of Likert data. The continuous ranked probability score for a predictive distribution F and outcome y is

$$\text{CRPS}(F, y) = \mathbb{E}_F|X - y| - \frac{1}{2} \mathbb{E}_F|X - X'|,$$

and it is *strictly proper*: uniquely optimised by $F = P$. Equation (1) is its first term alone. Deleting the second is exactly what makes a constant optimal.

This yields the natural repair. Define

$$S_\lambda = 1 - [\mathbb{E}|X - y| - \lambda \mathbb{E}|X - X'|]/4, \quad (2)$$

$\lambda=0$ is the audited metric — improper, constant-optimal. $\lambda=1/2$ is CRPS — strictly proper, truth-optimal. $\lambda>1/2$ is improper in the opposite direction. $\lambda=1/2$ is the unique proper choice, not a tunable hyperparameter.

Sweeping λ over the same 20 cells (Table 1) shows the repair works and what it costs the simulator. At $\lambda=1/2$ the human ceiling (0.859) rises decisively above the constant floor (0.796): the ordering that §7 shows is inverted under the audited metric is *correct* under the proper one. Humans overtake the constant at $\lambda^*=0.30$, comfortably below the proper coefficient. The model needs $\lambda^*=0.75$ (median grid crossover; 0.757 on the aggregate curve) — a rule that over-rewards dispersion — and beats the constant in only 2 of 20 cells at the proper $\lambda=1/2$, though it does overtake in most cells once λ reaches 1, where the reward for dispersion is twice proper. So the simulator’s failure is not an artefact of a bad metric: it survives the correct one.

so larger is better. Scores in the appendices are reported as the corresponding *loss* $L_{1/2} = 1 - S_{1/2}$, the normalised CRPS, where smaller is better; we use the loss wherever a table compares errors and mark the direction in each caption.

Table 2 evaluates the bound on the real data: the structural penalty against a faithful simulator averages **7.75 points**, and $\text{GMD}_j \geq \text{MAD}_j$ holds on 120/120 items in all four clusters.

4 Empirical audit

Setup. We re-score an existing cluster-level digital-twin system on its own artefacts: the same 410,168-respondent IPIP-NEO-120 corpus, the same $k=4$ clusters, and the exact train/test respondent IDs it recorded, for five LLMs (GigaChat3-10B, GPT-4.1-mini, GPT-4.1-nano, Grok-4.1-fast,

λ	Model	Constant	Human
0 (audited)	0.726	0.796	0.720
0.25	0.749	0.796	0.789
0.5 (CRPS, proper)	0.772	0.796	0.859
0.75	0.795	0.796	0.929
1.0	0.819	0.796	0.998

Table 1: The S_λ family of (2), averaged over 20 cells, all values on the corrected orientation. A constant wins at the audited $\lambda=0$; at the proper $\lambda=1/2$ a real cluster member wins, as it must. All-pairs scores, slightly below the paired ones of §4; the ordering is the same under either.

Clu.	n	MAD	GMD	S_{const}	S_{faith}
0	107,174	0.804	1.107	0.7990	0.7232
1	108,880	0.723	1.008	0.8192	0.7480
2	108,867	0.809	1.128	0.7978	0.7181
3	85,247	0.871	1.204	0.7822	0.6990
Mean		0.802	1.112	0.7996	0.7221

Table 2: Proposition 1 on 410,168 IPIP-NEO-120 respondents. $S(\text{const})$ is the best constant predictor; $S(\text{faithful})$ is a simulator that reproduces the cluster response distribution exactly. The faithful simulator is worse in every cluster.

Qwen3-235B) $\times$ four clusters. We add no LLM calls: model answers are read from the released run artefacts. Baselines are fitted on the recorded train IDs and scored on the recorded test IDs.

Baselines. *Cluster mean*: the per-item rounded mean of the cluster’s training respondents, one constant vector. *Cluster majority*: the per-item mode. *Global mean*: the per-item grand mean over all respondents, ignoring clusters. *Empirical sample*: sample each item independently from the cluster’s training response distribution. *Real human*: copy the answer vector of a randomly chosen training respondent from the same cluster.

Result. Table 3 reports S . The evolutionary persona averages 0.660 as released and 0.739 after the orientation correction of §6; the constant cluster mean averages 0.796. The LLM loses in 20 of 20 cells under either orientation, by 13.6 and 5.8 points respectively. What survives correction is *matched-cell* dominance: the model loses to its own cluster’s constant in every cell, while its best cell (0.807) now exceeds the weakest constant elsewhere (0.779), so the comparison only holds within a cluster. A per-item global mean — which never looks at the cluster at all — scores 0.774, still above the persona’s corrected 0.739.

Model	Cluster				Mean
	0	1	2	3	
GPT-4.1-mini	0.657	0.638	0.710	0.652	0.664
GPT-4.1-nano	0.673	0.678	0.697	0.670	0.679
GigaChat3-10B	0.711	0.691	0.737	0.672	0.703
Grok-4.1-fast	0.598	0.583	0.659	0.614	0.613
Qwen3-235B	0.634	0.615	0.677	0.642	0.642
<i>Cluster mean (constant)</i>	0.794	0.827	0.785	0.779	0.796
<i>Real human, same cluster</i>	0.720	0.747	0.715	0.704	0.721

Table 3: Answer similarity S on the original system’s own test IDs, in the *original* orientation; correction adds 0.079 on average and preserves all 20 matched constant wins, though it does reorder the models. The constant (bold) beats every LLM in every cell; a real human scores below it.

5 Does conditioning on the right cluster matter?

No persona study we know of reports the wrong-persona control, so no reported persona effect is currently attributable to the persona. We score every simulator against every cluster’s humans — 80 cells — and report diagonal dominance descriptively.

Metric	Matched	Mismatched	Δ
S	0.728	0.663	+0.065
DF	0.782	0.705	+0.077
C2ST	0.996	0.985	−0.010
VR	0.431	0.436	−0.005

Matched comparisons improve S and DF; they show no improvement in C2ST or VR. We report these gaps descriptively and attach no p -value: a permutation that shuffles the 80 cell labels freely destroys the matched structure of the five 4×4 matrices, and the models share clusters and personas, so its null is not the right one. Rescoring fixed outputs against different human populations also does not isolate a causal effect of changing the prompt. This is the behavioural counterpart of §8 — the persona moves which items get endorsed, and not who is endorsing them — and it makes the wrong-persona control as necessary a baseline as the constant floor.

6 An orientation defect, and what survives it

Before comparing anything, the two sides must be on the same scale. They are not.

The released corpus stores human answers *reverse-recoded* so that a higher number always means more of the trait, while the simulator is shown the literal item text — “Cheat to get ahead” — and asked how accurately it describes it, i.e. it answers *raw*. On the 55 negatively-keyed items of

the IPIP-NEO-120 the two scales run in opposite directions, and every published model-vs-human comparison is computed across that flip.

Two signatures make this unambiguous. Negatively-keyed items average 3.373 against 3.425 for positively-keyed ones — level, whereas raw self-reports on undesirable content sit far lower. And “Cheat to get ahead” carries a human mean of 4.15 out of 5.

The effect is large and it runs *against* the simulator. Correcting the orientation raises the correlation between the model’s and the humans’ per-item mean profiles from +0.09 — indistinguishable from no relationship at all — to +0.63, and raises S from 0.660 to 0.739.

Three conclusions survive, with their invariances:

- **Proposition 1 is invariant.** MAD and GMD are computed on the human distribution alone, and both are reflection-invariant.
- **The variance ratio is exactly invariant,** since $\text{sd}(6 - x) = \text{sd}(x)$. The under-dispersion finding of §8 is untouched.
- **The constant still wins every cell.** Reorientation closes the gap from 13.6 to 5.8 points, but the constant cluster mean beats the corrected model in 20 of 20 cells.

Correcting the defect in fact *sharpens* the paper’s claim. The corrected model scores 0.739, *above* the 0.721 of a real member of the target cluster — while a classifier separates that same model from real members at AUC 0.994 (§8). A metric that ranks a near-perfectly detectable simulator above the humans it fails to imitate is not measuring human-likeness under any reading.

We recommend an orientation check as a standing first step for benchmarks of this kind: correlate




Figure 1: Corrected orientation, 20 cells. Left: under the old metric the constant outscores a real human. Right: under C2ST-AUC (lower is better) the human sits near chance (0.5) while model and constant are both highly separable. Only the right panel places the human above the constant.

the simulator’s per-item mean profile against the population’s, and treat a near-zero correlation as a scale defect until shown otherwise.

7 The human ceiling sits below the constant floor

The decisive check is what a *real* member of the target cluster scores. A cluster-level digital twin cannot do better than being an actual member of the cluster; that is the definition of success.

A randomly drawn same-cluster respondent scores 0.721 — and is beaten by the constant cluster mean in 20 of 20 cells (Table 3, Figure 1, left). This is exactly Proposition 1 realised: a real human *is* a draw from P_j , so their expected score is $1 - \text{GMD}/4$, which the constant’s $1 - \text{MAD}/4$ must exceed. The predicted values (0.7996 constant, 0.7221 faithful) match the measured ones (0.796, 0.721) to within 0.4 points.

A measure that ranks a fixed vector above a genuine group member is not measuring group membership. Any ranking it induces over simulators — including the ranking prior work uses to select prompts — is therefore uninterpretable as evidence about personality fidelity.

8 A replacement suite

A usable metric must (i) be maximised by distributional fidelity rather than central tendency, and (ii) be reported against an explicit *human ceiling* and *constant floor*, so a number can be read as progress toward a real target rather than in the abstract. We propose three, computed from a set of simulated respondents rather than from paired individuals.

C2ST-AUC (primary). Train a gradient-boosted classifier to separate n simulated answer vec-

tors from n held-out human vectors of the same cluster; report cross-validated AUC, folded to $\max(\text{AUC}, 1 - \text{AUC})$. An AUC near 0.5 means this classifier cannot separate the samples at this sample size — evidence of indistinguishability under the stated configuration, not of distributional equality (App. N). Unlike S it cannot be gamed by collapsing onto the mean.

Distributional fidelity. $1 - \frac{1}{J} \sum_j W_1(\hat{P}_j^{\text{model}}, \hat{P}_j^{\text{human}})/4$, with W_1 the per-item Wasserstein-1 distance. Rewards matching the full response distribution.

Variance ratio. $\frac{1}{J} \sum_j \text{sd}(\text{model}_j)/\text{sd}(\text{human}_j)$; 1.0 is ideal, a constant gives 0. A diagnostic, not an optimisation target.

Results. Table 4 applies the suite to the same 20 cells. All three order human > constant, as required; the old metric does the reverse. Two substantive findings follow. First, current personas are *near-perfectly separable* from real humans (C2ST-AUC 0.994 against a human ceiling of 0.564): a classifier separates a persona from the group it simulates almost perfectly, while barely separating one human sample from another. The constant scores 0.926 here, but we do not read 0.994 against 0.926 as an intrinsic ordering — constant detection reaches 1.000 with more data, so the gap between two already-detectable systems is a property of the classifier. Second, they are systematically under-dispersed: $\text{VR} = 0.43$, a ratio of standard deviations, so LLM personas produce well under half the response *spread* of the humans they simulate (Figure 2). Evolutionary optimisation moves none of these appreciably.

Under-dispersion is the wrong summary. A single variance ratio hides which variance is missing. Writing each answer matrix as $Y_{pj} = \mu + a_p + b_j + r_{pj}$ — respondent main effect, item main effect, and the two-way residual — tells a sharper story. One observation per person–item cell cannot separate a stable interaction from response noise, so we read this term as a residual, not as individuation. The simulator emits nearly the *right total* variance (ratio 0.942 over the 20 cells) but places it differently: the item main effect takes 81.2% of its variance against 29.8% for humans, while the residual takes 17.7% against 67.0%. The residual SD ratio is 0.466, and 0 for a constant by construction.

So the simulator is not emitting too little. It is emitting a confident cluster-level *answer key* —

Metric	LLM	Human ceiling	Const. floor
C2ST-AUC ↓	0.994	0.564	0.926
DF ↑	0.783	0.939	0.795
VR (1.0 ideal)	0.431	1.018	0.000
$S_{\text{answer}} \uparrow$ (old)	0.739	0.721	0.796 †

Table 4: The suite over 20 cluster $\times$ model cells, each metric against its human ceiling and constant floor. †: under the old metric the constant floor *exceeds* the human ceiling — the failure this paper documents.




Figure 2: Response *spread* relative to the humans being simulated, as a ratio of standard deviations. Every model is substantially under-dispersed.

which items are endorsable — with little variation in which *person* endorses which item. We recommend reporting the residual ratio alongside the plain variance ratio, since the two rank systems differently. Individuation itself we measure with the aligned amplitude of §11 and the correct-versus-permuted CRPS, which noise cannot inflate: permutation leaves the residual unchanged.

8.1 What is C2ST actually reading?

A metric that separates simulators from humans at AUC 0.994 invites the objection that the classifier detects formatting rather than personality. A feature ladder (App. I) shows the objection has real force: five content-free response-style scalars already recover most of the separability above chance. It is not only style — ipsatising each respondent, or using facet scores alone, still separates nearly as well — so style and content are redundant routes to the same conclusion. We therefore recommend reporting C2ST with its style-only rung alongside it, so readers can see how much of any separability is carried by response style. A simulator that closed only the style gap would look far better on the headline number while remaining just as distinguishable on content.

The anchors are estimated at $n=40$ per side. Ap-

pendix N sweeps n and shows the C2ST human ceiling is entirely finite-sample classifier bias (extrapolating to 0.500) while the DF ceiling is entirely plug-in bias (extrapolating to 1.000, i.e. pessimistic). Below $n \approx 40$ our classifier configuration stops separating at all, so we report every anchor with the n and the settings it was measured at rather than proposing a universal floor.

9 The same penalty, in a fitted decoder

The proposition concerns a faithful sampler against a constant. A fitted predictor shows the same pressure. Fitting on half of each cluster and scoring on the other half, two oracles differ only in what they emit: a **mean** decoder (ridge on the ordinal code) predicting the conditional mean, and a **residual-resampled** one, the same fit plus a resampled residual, which restores marginal spread. The mean decoder scores $S = 0.802$, the residual-resampled one $S = 0.730$ — worse by 7.2 points for being closer to the response distribution on that margin. This is an illustration, not a test of Prop. 1: the decoder conditions on a richer code than the belief prompts carry, so its 7.2 points are not a prediction of the GMD — MAD gap. It beats the constant (0.792) by 0.9 points — what one fitted reader achieved, not a bound (App. K).

10 Does this generalise beyond one questionnaire?

Every result so far comes from IPIP-NEO-120, so the objection is that the audit describes that questionnaire rather than the task. We re-test on an independent instrument, construct and population: the Short Dark Triad (SD3), 17,740 respondents on 27 items.

The structural findings replicate: $\text{GMD} \geq \text{MAD}$ on 27/27 items with a penalty of 8.67 points against IPIP’s 7.75, and the constant again outscores a real respondent (0.744 vs 0.664) in 200 of 200 bootstraps. The defect belongs to the metric applied to ordinal questionnaire data, not to the Big Five or one corpus.

Both behavioural findings of §11 transfer to SD3, and under a design where the prompt cannot contain the answers being scored; on IPIP that design weakens one of them (App. B).

11 Where the variance is lost

Knowing the simulator is mis-dispersed does not say *where* the dispersion goes. We read the model’s

own belief before any sampling occurs, on open-weight models served locally.

Two stages. For each persona and item we take the model’s next-token distribution over the answer tokens $\{1, \dots, 5\}$ directly from the logprobs, then sample each item from its own belief and score the resulting population. One implementation detail matters: hybrid reasoning models emit a `<think>` token first — with probability ≈ 1 — so a naive readout sees no mass on the scale at all; we disable thinking and prefill the assistant turn.

Result. What is collapsed is the *individuating* part of the belief, not the belief’s total spread. In the 2024–25 models both are small; in the 2026 reasoning-first models the total spread reaches or exceeds the human marginal while individuation stays as low as ever.

The statistic matters. The quantity comparable to a human *across-respondent* SD is not the average within-persona conditional SD but the SD of the persona-*mixture* distribution, from the law of total variance $\text{Var}(X_j) = \mathbb{E}_i[\text{Var}(X_j | i)] + \text{Var}_i(\mathbb{E}[X_j | i])$, computed exactly from the retained probability vectors. The within-persona term alone would understate a faithful simulator, which can hold sharp per-persona beliefs and still match human variance through differing persona means.

Model	Params	VR _{with}	VR _{betw}	VR _{tot}
Qwen2.5-7B-Instruct	7B	0.367	0.193	0.430
Qwen2.5-32B-Instruct	32B	0.242	0.220	0.343
Qwen2.5-72B-Instruct	72B	0.367	0.295	0.488
Mistral-Small-24B	24B	0.538	0.177	0.579
QwQ-32B (reasoning)	32B	0.729	0.154	0.753
Qwen3-235B-A22B-2507	235B	0.695	0.258	0.758
Qwen3.5-9B	9B	1.058	0.167	1.078
GLM-4.7-Flash (MoE)	29B	1.209	0.265	1.242
gpt-oss-20b (MoE)	21B	1.121	0.126	1.131
Granite-4.1-8B	8B	0.386	0.207	0.457
Apertus-8B-Instruct	8B	0.927	0.130	0.939
<i>human</i>	—	—	—	1.000

Essentially all probability mass lands on the five valid answer tokens (0.97–1.00), so no model has difficulty *emitting* a number. Belief entropy nonetheless spans 0.42–1.38 nats (maximum $\ln 5 = 1.609$), and we can offer no clean account of it: it tracks neither release date (Granite-4.1, 0.49) nor reasoning-tuning (QwQ 0.97; the standard-instruct Apertus 1.24). The 235B model is the one the audited system used, so this measures that system, not a proxy.

Total dispersion and individuation are decoupled. The decomposition says something the total does not, and the evidence for this is existential

rather than correlational. gpt-oss-20b and Qwen2.5-72B invert cleanly: the former has $2.3\times$ the total dispersion of the latter (1.131 vs 0.488) and *less than half* its individuation (0.126 vs 0.295). The rank correlation over the eleven models is $\rho = -0.30$ with a bootstrap 95% interval of $[-0.92, 0.42]$.

QwQ-32B sharpens it: its dispersion is almost all per-item uncertainty (0.729 of 0.753), and it individuates less (0.154) than far smaller models.

The Qwen2.5 family isolates this at fixed family and varying scale: individuation rises monotonically, $0.193 \rightarrow 0.220 \rightarrow 0.295$ at 7B, 32B, 72B, while the total does not, moving $0.430 \rightarrow 0.343 \rightarrow 0.488$.

One bound survives across the whole set — a $34\times$ parameter range, six families, three generations, dense and mixture-of-experts, and both standard and reasoning-tuned models: no model here exceeds $\text{VR}_{\text{between}} \approx 0.30$ — a bound on this set, not a law: Gemma-4-12B, measured later, reaches 0.334 (App. L). Individuation stays under half the human value throughout, and neither the largest model nor the most dispersed one is the best at it.

The persona helps, and still loses. Two questions have run together: does the simulator reproduce the cluster’s distribution, and does it identify *this* respondent? Scoring each persona against that person’s own answers under normalised CRPS, and permuting belief rows among respondents within a cluster — which leaves every population statistic identical while destroying identity — separates them. All eleven score better on the correct assignment than the permuted one by 0.004–0.017, every paired interval excluding zero, yet all remain worse than a cluster empirical prior fitted on *other* respondents — which uses cluster membership but no persona — by 0.036–0.095. With conditioning and target items disjoint, the gain is resolved in three of four tested cells and unresolved for Granite-4.1 on IPIP, while all four still lose to the prior (App. B).

The two facts reconcile exactly, and without any fitted reference. Writing I for the identity gain above, $\bar{q}$ for the model’s within-cluster mean CDF and D_Q for the across-persona variance of its CDFs,

$$R(P_0) - R(Q) = I - D_Q - [R(\bar{q}) - R(P_0)], \quad (3)$$

an identity in measured quantities, exact on our data to machine precision. It locates the loss. The bracket — how much worse the model’s *mixture* is

than the prior — is 0.038–0.090, whereas the net identity contribution $D_Q - I$ lies within ± 0.005 for every model, in absolute value below a tenth of that model’s mixture excess. Spread across personas can help or hurt slightly; what decides the comparison is that the mixture itself is wrong.

The prompt is not what caps individuation. A persona is individuated only through a five-word quantisation of ten numbers, but that channel carries more than the models extract: a supervised conditional-mean decoder fitted on held-out respondents reaches 0.482 on these same personas, where the best model delivers 0.295, and no two personas share a code (App. K). The decoder is a reference, not a bound — a mis-calibrated simulator could exceed it by dispersing wrongly — so it shows the conditioning carries signal.

How much of it points the right way. VR_{between} measures how far apart a simulator places its persona means, not whether that spread identifies the right respondent. Correlating each model’s persona means against the oracle’s over the forty personas — both on the corpus’s recoded scale — gives $r = 0.16$ – 0.52 . The aligned amplitude $A = \frac{1}{J} \sum_j r_j \text{sd}(M_j) / \text{sd}(H_j)$, in units of human item SD, is 0.033–0.183, against 0.498 [0.461, 0.517] for real respondents: the best reaches about a third of that, the weakest under a tenth. It separates models VR_{between} ranks together — GLM-4.7 and Qwen3-235B sit 0.007 apart in spread, yet differ in A by $+0.017$ [$+0.009, +0.025$] — and rises with scale across Qwen2.5 in three non-overlapping steps ($0.067 \rightarrow 0.116 \rightarrow 0.183$).

A simulator can therefore look more dispersed simply by being less confident: VR_{total} confounds uncertainty with individuation, and VR_{between} is what a persona claim should be judged on.

Total dispersion is not uniformly low: it is 0.758 for the largest 2024–25 model and exceeds the human marginal for three others (1.078, 1.131, 1.242). Nor is it lost at decoding: sampling retains almost all of it ($1.078 \rightarrow 1.061$, $1.243 \rightarrow 1.218$, $1.131 \rightarrow 1.115$). What stays limited at every scale is *individuation*.

For the low-dispersion models the spread is likewise not discarded at decoding; it was never present, so recovering it there means *adding* dispersion, not revealing it: $p^{1/T}$ lifts Qwen2.5-7B from 0.43 to 1.02 at $T=5$ while supplying none of the alignment A measures. And since all mass lands on valid answer tokens, accounts blaming

mode collapse on difficulty *emitting* numbers are unsupported here.

12 Is $k=4$ a partition at all?

The audited systems condition on a $k=4$ k -means partition of the 30 facet scores. We reproduce it ($ARI = 0.946$) and test $k \in \{2, \dots, 15\}$ under three preprocessings: no criterion supports $k=4$, and the bootstrap stability that appears to defend it is matched exactly by structureless surrogates (App. J). The clusters are a convenient partition, not a discovered structure — which is why our results are stated for the partition the systems use rather than for personality types.

13 Related work

Personality conditioning of LLMs is commonly validated by administering psychometric inventories to the conditioned model (Serapio-García et al., 2023; Jiang et al., 2024), with well-documented prompt sensitivity (Han et al., 2026; Gupta et al., 2024). Individual-level simulation from transcripts achieves high behavioural fidelity (Park et al., 2024) at a per-person cost that motivates the cluster-level shortcut we study; evolutionary prompt search (Guo et al., 2024; Prasad et al., 2023; Zhou et al., 2022) supplies the optimiser. Our contribution is orthogonal: the objective they are scored against is maximised by a degenerate solution.

Classifier two-sample tests are standard for comparing distributions (Lopez-Paz and Oquab, 2017); the failure mode we identify, an absolute-error objective minimised by a conditional central tendency, is why such losses blur predictions in generative modelling.

14 Conclusion

Cluster-level LLM personality simulation is evaluated with a metric that a constant vector wins and a real human loses. We proved why, measured the penalty at 7.75 points, confirmed the inversion in 20 of 20 matched cells, and proposed a suite that orders humans above constants. Beneath it the persona helps all eleven models and none enough, the loss sitting in the mixture rather than identity. Across four instruments improper scores select measurably worse simulators; score with $S_{1/2}$ and report floor and ceiling.

Limitations

Our audit covers one corpus (IPIP-NEO-120), one clustering pipeline, one optimisation family, and five models; the belief-level measurement of §11 covers eleven open-weight models spanning 7B–235B, six families and three generations, and the cross-instrument replication of §10 carries the structural findings and the identity result but not the full suite, which needs sampled simulator outputs where SD3 has only the two belief readouts of App. B; belief-level measurement anywhere needs logprob access the API models in the audit do not provide, and three current large mixture-of-experts models — GLM-5.3-Flash, Qwen3.8-Flash-Next and DeepSeek-V4-Flash — could not be read at all on this protocol, returning NaN or zero mass on the answer tokens or shipping no chat template, so the candidate pools here reach 35B and not beyond; individual readouts are reproducible bit-for-bit across clusters and software stacks but *not* across accelerator architectures, where 7.7% of modal answers flip, so the per-item beliefs we release should be read as machine-specific even though the aggregates they feed move by at most 0.016 (App. F); the within-family trend rests on three points and the cross-family comparison on eleven, so both are directional rather than fitted laws; all models are read with thinking disabled and the assistant turn prefilled, and the four least confident of them (belief entropy above 1 nat) should not have their above-human total dispersion read as calibrated uncertainty, though low confidence cannot inflate VR_{between} , which is computed from persona means; two elicitation details proved to matter more than we expected — reading these models through a chat-completions server returned almost no mass on the answer tokens where the forward pass recovers 0.99–1.00, and harmony-format models require the assistant turn to open a channel before the pre-fill, without which mass on the scale falls to 0.39 — so a null result in this literature may report a readout rather than a model; precision varies across the models we report — nine run at bf16, gpt-oss-20b ships MXFP4 and is dequantised to bf16 at load, and Qwen3-235B is an FP8 checkpoint served at FP8 — and a further FP8 checkpoint resisted measurement through both available routes, so we make no claim that individuation is invariant to quantisation; the proposition is general to the absolute-error metric in (1), but the empirical magnitudes are specific to this setting. The model answers we re-score

were produced by prior runs we did not control, at temperatures we did not set, and one released run was committed with entirely missing answers and is excluded. Our human ceiling is a single same-cluster respondent, which underestimates what a well-specified individual-level twin could achieve; we make no claim about individual-level methods. C2ST-AUC depends on classifier capacity and on $n=40$ per side, so its absolute value should be read as a lower bound on separability, and §8.1 shows most of what it reads is response style; we fold the AUC as $\max(\text{AUC}, 1 - \text{AUC})$, which places the null above 0.5 (0.517 signed against 0.569 folded) and makes our small- n floor a property of this classifier rather than of C2ST; and DF treats items independently and so ignores within-respondent answer dependence. The corpus is self-reported, skewed toward respondents aged 18–25 and toward the USA, UK and Canada, so neither the clusters nor the ceilings should be read as representative of any broader population. Finally, we show current personas fail this evaluation; we do not show that any method can pass it.

Ethics statement

We use a public, de-identified corpus of self-reported personality questionnaire responses and add no new human-subjects data collection. Our results are cautionary: they indicate that current cluster-level personality simulations are readily distinguishable from real respondents and under-represent human response variance, and so should not be used as substitutes for human participants in research, product, or policy decisions. Reporting an evaluation number without its constant floor risks overstating the fidelity of synthetic respondents, which is the specific harm this paper aims to reduce. Simulating demographic or personality groups also risks reinforcing stereotypes when cluster descriptions are treated as characterisations of real people rather than as statistical summaries.

A Reproducibility

Every analysis in this paper is deterministic given its inputs and requires no LLM calls except the belief-level measurement of §11. The released corpus and committed run artefacts reconstruct the scoring, selection and orientation analyses. Some inputs are too large for the repository and are not in it: the refreshed belief tensors behind App. L, the per-respondent score arrays behind App. M,

the $n=512$ forecast tensors behind App. D, and the cross-cluster and cross-architecture replication readouts behind App. F. These accompany the submission as one archive (568 files, 92 MB compressed) carrying a per-file SHA-256 manifest and the CPU commands that regenerate each number above from it; no GPU and no model weights are required, because the archive holds the saved model outputs rather than the runs that produced them. We verified that path end to end: from the archive alone, the intervals of App. L, the frequencies of App. M, the eight paired differences of App. D and the deviations of App. F all reproduce as printed. We do not claim the list is exhaustive; where an analysis consumes a tensor, the tensor is in the repository only if this appendix says so, and otherwise in the archive. We ran the full analysis pipeline independently on two clusters with different operating systems, Python versions and library versions (scikit-learn 1.2 vs 1.7, pandas 2.1 vs 3.0). Every reported figure we compared — the λ^* values, the orientation deltas, the variance components, and the null-calibration ARIs of App. J — agreed to the digits we print, except the $k=4$ reproduction ARI, which is 0.946 on one environment and 0.951 on the other. Code and per-cell results are released. The belief simplices ($40 \text{ personas} \times 120 \text{ items} \times 5$) per cluster are included for the eleven models of §11, with the persona ordering fixed by the released corpus. Runs that failed the validity guard ship too, marked invalid, so the record shows what was measured and rejected. One transform maps model answers onto the corpus’s recoded scale and is imported by every analysis that compares the two.

B A disjoint conditioning/target replication

Every belief result in the body conditions on scores computed from the same items it is then scored on, so a sceptic can say the models reconstruct arithmetic rather than predict a person. We remove that by construction. Items are split within each scale, the persona is built only from the input half, and the readout covers only the target half: two input and two target items per IPIP facet (60/60), four input and five target per SD3 trait (12/15). No cluster label is used — the released IPIP clustering is derived from all 120 items and would reimport target information — and the prior and the decoder are fitted on training respondents whose target answers are never seen. The five score bins are fixed equal-

width bands of the 0–100 range, not fitted. This design differs from the body’s in more than the split: the persona is one generated line per facet rather than the body’s per-cluster prose, respondents are sampled globally rather than within clusters, and only half the items are scored. It is a replication under a stricter design, not a controlled ablation of overlap.

On 256 held-out respondents the second fact of §11 is unqualified and the first is weakened. Every model still loses to the empirical prior, by 0.098–0.204 across all four model–instrument cells. The identity effect survives in three of those four: on SD3 by $+0.010$ [$+0.006, +0.016$] for Qwen2.5-7B and $+0.017$ [$+0.013, +0.020$] for Granite-4.1, and on IPIP by $+0.004$ [$+0.002, +0.005$] for Qwen2.5-7B — but Granite-4.1 on IPIP gives $+0.002$ [$-0.001, +0.005$], which does not exclude zero.

One thing follows and one does not. Positive identity gains occur with conditioning and target items disjoint, in three of four cells, so item overlap is not necessary for them. What this design cannot do is apportion the body’s IPIP gain between overlap and identity, because it changes the persona text, the respondent sample and the scored items as well as the split: Qwen2.5-7B’s 0.006 there and 0.004 here are not two measurements of one quantity, and an interval excluding zero in one design beside one that does not in another is no evidence of a decrease. What survives every cell is that the simulator loses to a prior that ignores the persona.

C Does the replacement change a decision?

An audit that only condemns a metric leaves the practical question open: if you adopt the replacement, do you ship something different, and is it better? We fix one candidate set and select twice.

Candidates are the eleven belief models of §11, re-run under the disjoint-item prompts of App. B; the body’s readouts cannot serve, since they condition on scores computed from the answers being scored. Qwen3-235B is excluded: its FP8 checkpoint read correctly on one allocation and returned near-zero mass on the answer tokens on others, so it fails the same ≥ 0.5 mass criterion we apply throughout, and reading it through a different execution path would give one candidate an elicitation route the other ten did not have. Both rules see the same probability forecasts, never greedy

answers: rule A ranks by expected answer similarity, the audited S_0 ; rule B by $S_{1/2}$. Selection uses 128 respondents; evaluation uses 128 more, disjoint from those, from the 20,000 training respondents on which both baselines are fitted, and from App. B’s panel. Only the selected models are run on the evaluation panel.

The rules disagree. Rule A selects Qwen2.5-72B ($S_0 = 0.753$); rule B selects Qwen3.5-9B ($L_{1/2} = 0.187$) — models an order of magnitude apart in size. On the untouched panel rule B’s pick is better by 0.010 [0.004, 0.016] under $L_{1/2}$ (lower is better), paired over respondents. The task is ordinal prediction, so that is the outcome we report; categorical log-loss agrees and is secondary, at 1.37 nats [1.26, 1.48]. The log-loss gap is the larger number but the weaker argument, since it mostly registers the 72B’s overconfidence rather than its ordinal accuracy — its errors are confident, not distant. Its one advantage is independence: neither selection rule uses log-loss, whereas $S_{1/2}$ is rule B’s own criterion, evaluated on disjoint respondents.

Two caveats belong with this. It is one candidate set, one split and one instrument, so it shows the evaluation changed a decision for the better — not that it reliably does. And the improvement is relative: both picks lose to the training-fitted empirical prior (0.155, 1.41 nats), so the replacement recommends a smaller model and simultaneously reports that no candidate beats ignoring the persona. Adding the degenerate candidates makes the point sharper still. With the constant vector in the pool, rule A selects *it* over every model ($S_0 = 0.766$ against the 72B’s 0.753), which is Prop. 1 arriving as a deployment recommendation; rule B selects the prior. The constant’s log-loss is infinite, and we report it as such rather than clipping it.

D Which scoring rule selects best?

App. C compares two rules on one instrument, and says plainly that this cannot show the replacement *reliably* chooses better. Here we widen it: nine scoring rules, twelve open-weight models spanning 0.5B to 35B across six families, and four instruments, two of which the body never uses.

The rules span three classes. Strictly proper: $S_{1/2}$ (ranked probability score), Brier, log-loss, spherical. Weakly proper: RMSE of the forecast mean, and mode accuracy. Improper: the audited S_0 , MAE of the forecast mean, and Wasserstein

distance to the observed point mass, which coincides with S_0 by construction here and so is not independent evidence. Instruments are IPIP-NEO-120, SD3 and IPIP-FFM-50 on five points, and HEXACO-PI-R on *seven* — the last is the only place the pipeline is exercised at $K \neq 5$. The first three are keyed from published item keys; the HEXACO distribution we use ships none, so its reverse-keyed items are inferred from each item’s correlation with the rest of its scale, and HEXACO results carry that caveat. Each rule selects on 512 respondents and is graded on 512 disjoint others, with baselines fitted on a further 8,000; all 96 read-outs put at least 0.92 of their mass on the answer tokens.

We score a selector by the *regret* its pick incurs, normalised within the candidate set so 0 is the best of these twelve models and 1 the worst, and averaged over the strictly proper evaluators *excluding the selector’s own rule*, so no rule is graded on the criterion it optimised. Because eight rules are compared against one reference, we report Holm-adjusted p and Bonferroni simultaneous intervals, and call a difference resolved only when both agree. The four differences below with intervals excluding zero all carry Holm $p \leq 0.01$; the three strictly proper alternatives carry $p = 1.00$.

rule	class	regret vs. $S_{1/2}$ (Bonferroni CI)
RMSE of mean	weak	+0.000 [+0.000, +0.000]
Brier	strict	+0.003 [−0.008, +0.025]
log-loss	strict	+0.004 [−0.007, +0.022]
spherical	strict	+0.007 [−0.002, +0.029]
MAE of mean	improper	+0.025 [+0.007, +0.055]
mode accuracy	weak	+0.238 [+0.084, +0.262]
S_0 audited	improper	+0.298 [+0.250, +0.368]
Wasserstein	improper	+0.298 [+0.250, +0.368]

Each row grades the two selectors on the same evaluators — the strictly proper rules minus both contestants’ own criteria — so a rule is never compared on a set containing its opponent’s criterion but not its own. Differences are bootstrap means over respondents, $B=4000$; intervals are Bonferroni simultaneous bounds holding jointly over all eight comparisons. Every improper rule selects worse than $S_{1/2}$, and the audited S_0 is the worst of them. The gap is not marginal: mean regret is 0.259 for S_0 against 0.003 for $S_{1/2}$. On SD3, S_0 selects a model whose regret is 0.98 — all but the worst of the twelve — while every strictly proper rule selects the best. An evaluator using the audited metric on this pool would ship very nearly the worst simulator available.

The number depends on the pool. With eight

956 closely matched instruction-tuned models the same
 957 comparison gives $+0.028$ [$+0.011, +0.059$]: still
 958 resolved, ten times smaller. Regret is normalised
 959 within the candidate set, so it measures position
 960 within that pool and cannot be read across pools.
 961 What it does show is where in *this* pool the audited
 962 metric lands: near the bottom on SD3, where its
 963 selection carries regret 0.98, while the proper rules
 964 select the best.

965 Three things this does not say. The strictly
 966 proper rules remain mutually indistinguishable, so
 967 the evidence is for the class boundary, not a rank-
 968 ing inside it. RMSE of the forecast mean — only
 969 weakly proper — is not distinguishable from $S_{1/2}$
 970 at all: it selects the same model on every instru-
 971 ment, and its paired difference is exactly zero. Elic-
 972 iting the mean is sufficient on this candidate set,
 973 and we have no basis for preferring a strict rule
 974 over it. And regret is normalised within the candi-
 975 date set, so 0 means the best of these twelve, never
 976 good: every candidate still loses to the persona-free
 977 prior of App. B.

978 Mode accuracy is the other collapse, at $+0.238$.
 979 Discarding the distribution and keeping the argmax
 980 is as damaging as the metric this paper audits. If
 981 a practitioner takes one operational rule from this
 982 work, it is that a simulator must be scored on the
 983 distribution it induces, not on the answer it would
 984 most likely give.

985 Sample size decided the weaker version of this.
 986 At 128 respondents per panel with eight candidates
 987 the S_0 comparison was $+0.070$ [$-0.017, +0.273$]
 988 — unresolved. Quadrupling the panels resolved it;
 989 widening the pool made it large.

990 E When does impropriety actually bite?

991 App. D shows S_0 selecting all but the worst sim-
 992 ulator from a heterogeneous pool. It would be
 993 convenient to conclude that optimising S_0 is there-
 994 fore harmful wherever it appears. We tested that
 995 directly and it is false.

996 The system this paper audits was built by evolv-
 997 ing cluster-level system prompts to maximise S_0 .
 998 We take those evolved genotypes and the base geno-
 999 types the search started from, assemble both with
 1000 the original code, and read the same belief distri-
 1001 bution under each — scoring with the objective
 1002 that was optimised and with two proper rules. The
 1003 prediction we registered before running was that
 1004 S_0 would improve while the proper scores would
 1005 not.

model	ΔS_0	$\Delta L_{1/2}$	$\Delta \log\text{-loss}$
Qwen3.8-27B	+0.060	-0.039	-0.235
granite-4.2-8B	+0.046	-0.046	-0.630
gpt-oss-20b	+0.052	-0.043	-0.251
gemma-4-12B-it	+0.070	-0.068	-0.619

1006 Evolved minus base, 256 held-out respondents
 1007 over four clusters, paired bootstrap. Every interval
 1008 excludes zero and *every* direction is an improve-
 1009 ment: $L_{1/2}$ and log-loss are losses, so the negative
 1010 signs are gains. The prediction is falsified on all
 1011 four models. Across respondents the change in S_0
 1012 and the change in $S_{1/2}$ correlate at $r = 0.96$ – 0.99 ,
 1013 and 85–97% of respondents improve on both at
 1014 once.
 1015

1016 The two results are not in tension. S_0 is max-
 1017 imised by a constant vector (Prop. 1) — a *point*
 1018 *mass*, not a flat distribution — so its optimum is
 1019 under-dispersed. Writing $\bar{\sigma}$ for a model’s mean
 1020 predictive standard deviation, the Spearman corre-
 1021 lation between $S_{1/2}$ as a reward and $\bar{\sigma}$ is positive on
 1022 all four instruments ($+0.71$ to $+0.86$): the proper
 1023 score consistently prefers the more dispersed can-
 1024 didate. S_0 does not, and its correlation varies —
 1025 $+0.03$ on IPIP, -0.09 on IPIP-FFM, -0.40 on
 1026 HEXACO and -0.97 on SD3 — with the failure
 1027 tracking the penalty. On SD3, where the correlation
 1028 is -0.97 , S_0 ships a model with $\bar{\sigma} = 0.48$ against
 1029 the proper rule’s 1.10 and incurs regret 0.98; on
 1030 IPIP and IPIP-FFM, where it is near zero, both
 1031 rules select the same model and the regret is zero.
 1032 These are correlations over twelve candidates per
 1033 instrument and we read them as descriptive, not as
 1034 an identified mechanism.
 1035

1036 What we can say about optimisation is narrower.
 1037 Evolutionary prompt search improved both S_0 and
 1038 the proper scores on all four models tested, so opti-
 1039 mising an improper objective did not degrade the
 1040 artefact here. We have since measured the disper-
 1041 sion of the candidates a prompt search can visit
 1042 (App. M): over twelve persona conditions the rules
 1043 disagree on two models of four, and on one of those
 1044 the disagreement is stable under resampling, with
 1045 S_0 taking the less dispersed candidate. The degener-
 1046 ate direction is therefore reachable by prompt
 1047 search in this candidate set, at a small observed
 1048 cost. Model selection over a heterogeneous pool
 1049 is the other case we measured: an over-confident
 1050 weak model is already most of the way to a point
 1051 mass, so S_0 rewards it and ships it.
 1052

1053 The pattern across the two settings is narrower
 than “never optimise S_0 ”. In the candidate sets
 we measured, S_0 disagreed with a proper rule far

more visibly in the model-selection pool, where it shipped a near-worst simulator, than in the prompt pool, where the selected candidates scored within 0.001 CRPS of each other. We do not convert that into a claim that one setting is more dangerous than the other: the two are not measured on a common scale (App. M). We report this because it bounds our own claim: the audit establishes that S_0 is improper and that it selects badly, not that every system built with it is worse for having used it.

F Is a belief readout a measurement?

Every number in App. D treats a readout as a stable property of a model. That is an assumption, and this paper’s own Limitations record the readout *path* moving mass on the answer tokens from 0.39 to 0.99, so we tested it rather than appealing to determinism. Two models and two instruments were re-read with the same script, seed and respondent panels on independent systems.

Same architecture and dtype, different everything else. Between two clusters running different software — torch 2.11.0 with transformers 5.17.0 against torch 2.13.0 with transformers 5.16.1, both on H200 and both on the `bf16` path — the readouts are *bit-identical*: maximum deviation 0 over the full belief simplex and 0 of 76,800 item–respondent pairs changing their modal answer. This is a statement about two small models on one numeric path, and §F.1 below shows it does not extend to a change of dtype.

Different architecture. Against A100 (`sm_80`) instead of H200 (`sm_90`), on the same `bf16` path, they are not. Mean absolute deviation per probability is 0.008–0.029, and 5,946 of 76,800 modal answers flip — 7.7%. The divergence is not a mismatched comparison: the four cells score identical respondents, and the flips sit precisely on near-ties. Where a modal answer flips, the median gap between the model’s top two levels is 0.051 against 0.507 where it does not; for one cell that median gap is 0.000, exact ties resolved differently by different kernels.

What survives. The scores largely do. Averaged over the four cells, the change is 1.6×10^{-3} for S_0 , 3.4×10^{-3} for mode accuracy and 4.7×10^{-3} for $S_{1/2}$; the largest single-cell change is 1.6×10^{-2} , for $S_{1/2}$ on gemma-4-E2B/SD3. We do not convert these into a ratio against the regret differences of

App. D: raw score changes and candidate-range-normalised regrets are different units and the comparison would be meaningless. What we can say is bounded and direct — a change of architecture moves a model’s score by at most 0.016 here, while it moves 7.7% of that model’s individual answers.

We had expected the distribution-based scores to be the *stable* ones, which would have given §8 a second argument independent of propriety. They are the least stable of the three, and we report that rather than the argument we went looking for. The practical reading is narrower: a per-item LLM “belief” is not reproducible across accelerators and should not be published as though it were, while an average over hundreds of items and respondents is stable enough to carry an effect an order of magnitude larger than the noise.

One arm is missing from *this* comparison, and it is not a minor one. V100 (`sm_70`) would have tested the `fp16` path; the build available when we ran these checks ships no Volta kernels, so the comparison above spans two Hopper-class and one Ampere-class device only. The `fp16` path is nevertheless the path on which part of §11’s table was measured, which we address next.

F.1 The table is not on a single numeric path

Auditing the run records for §11 showed that its rows were not all produced on one device class. `h4_variance_ladder_hf.py` selects `bf16` at `sm_80` and above and `fp16` below it, and four rows — Mistral-Small-24B, Qwen2.5-7B, Qwen2.5-32B and QwQ-32B — were read on V100 (`sm_70`) in `fp16`, while the remainder were read on H200 (`sm_90`) in `bf16`. The appendix above tests `bf16` against `bf16`; it never tested the path those four rows used.

We therefore re-read all four on `sm_90` in `bf16`, holding the model, the script, the panel (40 personas per cluster) and the respondents fixed, so that only the device class and dtype differ. The belief readout is deterministic given a model and a panel — the seed enters only the sampled answers — so these are not repeated draws:

Model	<code>sm_70 fp16</code>	<code>sm_90 bf16</code>	Δ
Qwen2.5-7B	0.430	0.427	−0.003
Qwen2.5-32B	0.343	0.340	−0.003
QwQ-32B	0.753	0.824	+0.070
Mistral-Small-24B	0.579	0.771	+0.192

Two of the four reproduce to 0.003 in VR_{total} . Two do not: QwQ moves 9% and Mistral 33%, in both cases because the model emits flatter answer

distributions on the newer path — Mistral’s mean belief entropy rises from 0.729 to 0.984 — which inflates the within-persona term. Panel size is not the explanation: on sm_90 Mistral moves 0.002 between 40 and 160 personas.

The individuation column is not affected. The quantity §11 actually rests on is VR_{between} , and it survives the change of path where the total does not. Re-reading ten of the eleven rows on sm_90 bfl6 at a common 160-persona panel (QwQ at 40, the largest panel on which it returns a readout) reproduces every published VR_{between} to within 0.0064, and the four fpl6-origin rows are no worse than the six bfl6-origin ones:

Model	pub.	re-read	Δ
Qwen2.5-7B [†]	0.193	0.187	−0.006
Qwen2.5-32B [†]	0.220	0.221	+0.001
Qwen2.5-72B	0.295	0.292	−0.003
Mistral-24B [†]	0.177	0.176	−0.001
QwQ-32B [†]	0.154	0.148	−0.006
gpt-oss-20b	0.126	0.123	−0.003
Granite-4.1-8B	0.207	0.203	−0.004
Apertus-8B	0.130	0.125	−0.005
GLM-4.7-Flash	0.265	0.263	−0.002
Qwen3.5-9B	0.167	0.163	−0.004

[†] published row measured on sm_70 fpl6.

Mistral is the clearest case: its VR_{total} moves 0.192 between paths while its VR_{between} moves 0.001. That is what the mechanism predicts. A change of numeric path alters how sharp each persona’s answer distribution is, which is a within-persona quantity; it barely moves where those personas’ *means* sit, which is what VR_{between} measures.

The practical consequence is narrower than the mixed path first suggests. The individuation claims of §11 and the ordering of models by VR_{between} stand: every published value reproduces within 0.0064 on one path. What does not transfer is the total-dispersion column and the within/between share derived from it, for the two models whose distributions sharpen. We leave the published table as measured and flag the affected statistic rather than restating numbers, and note that App. L is uniformly sm_90 bfl6 throughout.

One further caution for anyone reusing the script: at 160 personas the QwQ read returned NaN for every cell while exiting normally, and was caught only by the validity guard that rejects an all-missing run. The same read at 40 personas succeeded. We have not traced the cause; a partially rather than wholly degenerate version of that fault would not trip the guard.

G Split sensitivity of the baselines

Re-drawing which 100 respondents are sampled at a fixed 60/40 ratio moves the parameter-free cluster-mean baseline by 0.54–0.77 points within a cluster over 20 draws, and the four-cluster average by 0.23 (pooling clusters gives 1.33, mixing cluster differences into split variation). Prior work’s gain is +0.72 points on one fixed split. This is context for how far a split can move a score, not a test of that gain: variation in baseline *levels* does not bound a *paired* improvement.

H What is a persona worth, in respondents?

“It loses to a constant” invites the reply that the constant is a strawman. So we restate the finding in a unit that is harder to dismiss: how many real respondents you would have to sample before a plug-in estimate of the cluster’s response distribution matches what the persona gives you.

Drawing m real respondents, forming the per-item empirical distribution, and sampling a synthetic population from it yields distributional fidelity 0.723 at $m=1$, 0.797 at $m=2$, 0.915 at $m=21$ and 0.923 at $m=34$. The curve is evaluated at the same sizes as the audited scores — forty synthetic against forty held-out humans — because DF depends on sample size, so a curve built at any other size cannot calibrate them. Reading the audited systems off it:

System	DF	Resp.-equiv.
Evolutionary LLM persona	0.783	2
Constant cluster mean	0.795	2
Human ceiling	0.939	off-curve

An evolutionary-optimised persona is worth about **two** real respondents: two is the smallest evaluated m whose mean DF exceeds the persona’s 0.783, and the constant sits at the same place. This is a threshold on the mean curve, not an equivalence test. The human ceiling of 0.939 is not crossed anywhere on the evaluated grid, up to $m=500$ (0.938), so we do not convert it into respondent units — a statement about the grid we searched, not a proof that no m reaches it. This is the same result as §4 in different units, and it is the version we would put to a practitioner weighing synthetic respondents against a small pilot survey.

I The C2ST feature ladder

A metric that separates simulators from humans at AUC 0.994 invites an obvious objection: perhaps

the classifier is detecting formatting rather than anything about personality. We test our own metric with a feature ladder, each rung scored identically against the same anchors. Orientation is handled per feature family: style features describe how a respondent uses the scale, so both sides sit on the literal scale the item was asked on, while content features compare trait-oriented quantities. The full-item rung reproduces the suite’s 0.994.

Features	LLM	Human	Const.
Full 120 items	0.994	0.564	0.926
5 style scalars only	0.894	0.584	0.988
Items, ipsatised	0.989	0.566	0.986
30 facet scores only	0.991	0.565	0.908
Forward-keyed only	0.990	0.568	0.810
Reverse-keyed only	0.993	0.583	0.843

Five content-free scalars — per-respondent mean and standard deviation, extreme -response rate, longest identical run, and forward-versus-reverse consistency — attain 80% of the full above-chance AUC on their own. The objection has real force, and we state it plainly: **response style alone reproduces most of the separability**, which is a statement about what suffices, not a decomposition of what the full classifier uses.

It is not only style. Ipsatising each respondent (removing their own mean and standard deviation, i.e. the two largest style features) still leaves AUC 0.989, and facet scores alone give 0.991. Style and content are redundant routes to the same conclusion rather than one masking the other. The human ceiling is also stable at 0.56–0.58 across every rung, so the anchor does not move with the feature set.

J Cluster validity in full

Since stability-based model selection is common, the general point is worth stating: at large n , bootstrap agreement must be read against a covariance-matched null or it certifies nothing.

No *separation* criterion supports $k=4$. Silhouette falls overall from 0.134 at $k=2$ to 0.053 at $k=15$, with minor upward steps, taking 0.082 at $k=4$; Davies–Bouldin is also minimised at $k=2$; the gap statistic increases monotonically and selects no interior k . A silhouette below 0.1 indicates essentially no separated structure.

Stability looks at first like a defence: bootstrap ARI over 40 resamples is 0.977, 0.952, 0.932, 0.903 at $k=2, 3, 4, 5$ and falls away after (0.797, 0.581), so $k=4$ looks close to the largest reproducible k .

That reading does not survive a null calibration. Rebuilding the pipeline on structureless surrogates — a Gaussian matched to the real 30×30 facet covariance, and a copula matched to the covariance and every per-facet marginal — reproduces the stability number exactly: ARI 0.911 and 0.921 at silhouette 0.075 and 0.078, against the real data’s 0.910 at 0.082. A separated 4-component control confirms the test has power (silhouette 0.672, ARI 1.000).

At $n=410,168$ any smooth unimodal cloud with this covariance is partitioned reproducibly, so ARI ≈ 0.91 measures estimator variance, not multimodality, and neither criterion supports $k=4$. The clusters are not demographic artefacts either (Cramér’s V of 0.055–0.057 against sex, age band and country); they are a reproducible discretisation of a continuum, which bounds the cluster-conditional signal any persona can exploit. Since stability-based model selection is common, the general point is worth stating: at large n , bootstrap agreement must be read against a covariance-matched null or it certifies nothing.

K What the persona channel carries

A persona reaches the model only through `get_modifier_bisect`, which writes each score as one of five words on width-20 bins; each cluster’s genotype names five traits and five of the thirty facets, and every other sentence of the system prompt is shared. Ten words therefore carry the respondent’s identity, and two respondents sharing those ten bins would receive byte-identical prompts and so must receive identical beliefs.

We estimate what that channel permits. Within each cluster we fit $\mathbb{E}[X_j \mid \text{code}]$ on respondents *excluding* the forty personas, then evaluate on those forty, so the quantity is directly comparable to VR_{between} and is never fitted on its own evaluation set.

Conditioning available to the model	Dec. SD
Ten scores, five words (<i>audited system</i>)	0.482
same, with boosted trees	0.483
Ten scores, written as numbers	0.512
All 35 scores, five words each	0.707
All 35 scores, written as numbers	0.726
<i>best of the eleven</i>	0.295

Three things follow. The quantisation is not the binding constraint: writing numbers instead of words buys 0.03. The channel is not degenerate: all forty personas in all four clusters receive distinct codes, so no model is forced to answer two

personas identically. And a boosted-tree decoder, which can use interactions, lands in the same place (0.483 against 0.482); that is consistent with the additive fit being close to the conditional mean, though two similar SD ratios do not prove it.

The oracle is fitted on this population and so knows its code-to-response mapping, which a zero-shot model does not; we therefore do not claim that a model *ought* to reach 0.482. The comparison rules out one explanation only — that the conditioning carries too little information to individuate — and that is the explanation the flatness of VR_{between} across eleven models most invites.

Uncertainty in §11. Identity gain and the prior comparison use 2000 percentile-bootstrap draws over respondents, 40 per cluster with equal cluster weight. The permutation control is not sampled but computed exactly, as the expectation of every donor row against every respondent; the prior and decoder are fitted once on the excluded respondents and held fixed, so intervals cover respondent sampling, not fitting.

Uncertainty on A. Persona rows are resampled jointly across the model, the decoder and the actual respondents, so paired comparisons stay paired; the decoder fit is held fixed, so the intervals describe persona-sampling uncertainty rather than fit uncertainty. Over 2000 draws the human reference is 0.498 [0.461, 0.517] and the eleven models span 0.033 [0.028, 0.036] for Apertus-8B to 0.183 [0.168, 0.193] for Qwen2.5-72B. The intervals are informative enough to separate models that VR_{between} ranks together: Qwen2.5-32B and Granite-4.1 differ by 0.013 in spread but their amplitudes do not overlap, [0.105, 0.122] against [0.071, 0.089]. Comparisons are made on the paired draws, not by asking whether marginal intervals overlap — two overlapping marginals can still have a difference interval excluding zero, as GLM-4.7 and Qwen3-235B do (+0.017 [+0.009, +0.025]) — and Qwen2.5-32B exceeds Granite-4.1 by +0.034 [+0.028, +0.040] despite a 0.013 gap in spread.

Where the shortfall falls. Each cluster names five facets, so twenty items have their own facet stated and one hundred are covered only at trait level. The named items carry $1.51\times$ more predictable-from-code signal (an SD ratio) (reference 0.671 against 0.444), and 10 of 11 models individuate them more in absolute terms, Qwen3-

235B being the exception (0.224 named against 0.264 elsewhere). Measured against those references the ratio is *lower* where the prompt is most specific — 0.353 on the named items against 0.433 on the rest, in the same direction for 10 of 11 models — but we report this descriptively and draw no causal conclusion from it, for two reasons. A facet score is computed from its own four items, so $\mathbb{E}[X_j \mid \text{facet}]$ is close to an arithmetic identity that a model without the scoring key cannot invert, while traits have the same structure diluted over twenty-four items; the named-item reference is inflated by more than the trait-level one, which alone could produce the gap. And the eleven models share personas, references and prompt structure, several within one family, so they are not independent replicates and we do not attach a p -value to the count.

Widening the channel does not fix it. The obvious remedy is to send more: we re-ran three models with every other word of the prompt held fixed and the ten modifier phrases replaced by the numbers themselves. The effect is inconsistent in magnitude — VR_{between} rises for Granite-4.1 (0.207 $\rightarrow$ 0.246) and gpt-oss (0.126 $\rightarrow$ 0.146) but falls for Apertus (0.130 $\rightarrow$ 0.096) — and, using the same joint draws so the comparison is paired, every model moves resolvably: A shifts by +0.053 [+0.045, +0.059] for Granite-4.1 and +0.008 [+0.005, +0.012] for gpt-oss but −0.013 [−0.016, −0.009] for Apertus. The effect is therefore not noise; it is real and model-dependent in sign. It is also inconsistent in direction, the aligned amplitude moving 0.081 $\rightarrow$ 0.134 for Granite-4.1 and 0.039 $\rightarrow$ 0.047 for gpt-oss but 0.033 $\rightarrow$ 0.020 for Apertus. Granite gains on both counts — its correlation with the oracle rises from 0.285 to 0.453 — so for that model the richer encoding is a real improvement, while Apertus loses spread and direction together. On three models a richer encoding is therefore a plausible but unreliable remedy, and we report it as a pilot rather than a recommendation.

The alignment does not depend on the oracle. Because r is measured against a fitted target, an oracle missing structure the models track would understate it. It does not. Boosted trees on the same code change r by at most 0.012 and a ridge on the continuous scores by at most 0.034; an oracle given all thirty-five scores lowers r substantially, by up to 0.168, which is the expected direction since models only ever see the ten-score code. The ordering of

models is unchanged throughout. Two independent computations agree where theory demands it: the ceiling as a standard-deviation ratio is 0.482 and the oracle’s correlation with actual respondents is 0.491, which coincide for a conditional mean. The check that assumes no oracle at all — correlating each model’s persona means directly against the forty respondents’ own answers — gives 0.076–0.248 against the oracle’s 0.491, so models recover between a sixth and a half of the attainable person-level alignment.

L Current-generation models, with intervals

The eleven-model table in §11 predates three of the four families below; gpt-oss-20b is re-measured rather than new. We repeated the belief-level measurement on one current model from each, under the identical protocol (40 personas per cluster, four clusters, exact readout from the retained probability vectors), and added the interval the original table lacks. Some readouts are missing for Gemma-4 — 30 persona-item readouts of 19,200, one in 640, all in cluster 1 where they are 30 of 4,800 (coverage 0.99375); those cells are held out of the moments rather than summed to zero, which would enter as an impossible answer of 0 and inflate the between-persona term. Every other model here has complete coverage. Resampling is over personas — the exchangeable unit for a between-persona variance — with 2000 draws; because the protocol assigns the same personas to every model, each replicate applies one set of persona indices to all four, so the differences below are paired.

Model	Params	VR _{with}	VR _{betw} [95% CI]	VR _{tot}
Granite-4.2-8B	8B	0.729	0.155 [0.143, 0.160]	0.757
Gemma-4-12B-it	12B	0.336	0.334 [0.303, 0.345]	0.506
gpt-oss-20b (MoE)	21B	1.120	0.127 [0.118, 0.130]	1.130
Qwen3.8-27B	27B	0.869	0.219 [0.197, 0.229]	0.911
human	—	—	1.000	1.000

The measurement replicates. gpt-oss-20b appears in both tables from independent runs. The new run returns $VR_{\text{between}} = 0.127$ against 0.126, $VR_{\text{within}} = 1.120$ against 1.121, and $VR_{\text{total}} = 1.130$ against 1.131. We report this because it is our only same-model repeat under this variance protocol (App. F tests readout stability separately), and it bears on §11 rather than on the new models.

The 0.30 bound is not a law. Gemma-4-12B measures $VR_{\text{between}} = 0.334$ [0.303, 0.345], above the ≈ 0.30 figure that held across all eleven earlier

models — but the interval’s lower end sits essentially *at* that figure, so we read this as evidence that the ceiling is a property of the models we happened to measure rather than as a decisive refutation of it. The qualitative claim is unchanged: 0.334 is still about a third of the human value. We have scoped §11 accordingly rather than restating the bound as universal.

The two failure modes are separable. Gemma-4 and gpt-oss-20b fail in opposite directions, and the intervals separate them. gpt-oss-20b carries *above*-human total dispersion ($VR_{\text{total}} = 1.130$) of which almost none individuates: the between-persona share of belief variance is 0.016 [0.014, 0.018]. Gemma-4 is its mirror — a between-persona share of 0.504 [0.466, 0.524] but collapsed in magnitude ($VR_{\text{total}} = 0.506$). It is not alone in that: Qwen2.5-32B and Qwen2.5-72B in the eleven-model set have between-persona shares of 0.485 and 0.429. The paired difference, gpt-oss minus Gemma-4, is -0.488 [-0.507 , -0.451]. Near-human total spread is therefore not evidence of individuation, and individuation is not evidence of human-like spread; a single variance ratio cannot distinguish the two, which is the argument of §11 restated on models chosen after it was written.

What does not separate. Of the six paired comparisons of the between-persona share, five exclude zero and one does not: Granite-4.2 against Qwen3.8 is -0.008 [-0.017 , $+0.003$]. Their point estimates, 0.068 and 0.076, must not be read as an ordering. We note it because the same two models *do* separate on VR_{between} itself (-0.064 [-0.073 , -0.049]): which quantity is reported decides whether these two models are distinguishable, and the point estimates alone would not have revealed that.

M Is the degenerate direction reachable by prompt search?

App. E reports that prompt search improved S_0 and the proper scores together, and notes that this alone cannot establish whether the degenerate optimum is out of reach — we had not measured the dispersion of the candidates a search can actually visit. This appendix measures it.

Design. We cross six persona directives with the base and evolved genotypes, giving twelve candidates scored on the same 192 respondents. Three

directives are intended to push towards concentration (*decisive*: avoid middle-of-the-scale answers; *extreme*: answer at the ends; *consistent*: similar items receive the same number), one is intended to push towards spread (*varied*), one asks for the spread of a single real respondent (*calibrated*), and one is the unmodified prompt. Each candidate is scored under S_0 and under CRPS, and we record its mean predictive SD. The verdict is fixed in code before the run: REACHABLE iff the candidate S_0 selects is strictly less dispersed than the one the proper rule selects.

Stability. A strict inequality between two close candidates is not a result, so we resample respondents (2000 draws, paired — all conditions are scored on the same respondents) and *re-run the selection inside each replicate*, so the reported stability includes the selection uncertainty rather than freezing the point-estimate picks.

Model	S_0 selects	rules differ	REACH.
Granite-4.2-8B	evolved/varied	0%	0.0%
Gemma-4-12B-it	evolved/varied	0%	0.0%
Qwen3.8-27B	evolved/varied	77%	77.4%
gpt-oss-20b	evolved/calibrated	99%	99.2%

Two point-estimate disagreements; one of them stable. On Granite-4.2 and Gemma-4 both rules select the same candidate in every one of the 2000 replicates: nothing to separate. On Qwen3.8 the point estimate fires REACHABLE, but the proper rule’s own choice is unstable — it takes *calibrated* in 77% of replicates and *varied* in the rest — so we report it as inconclusive rather than as a second instance. On gpt-oss-20b the disagreement is stable: S_0 selects *evolved/calibrated* in 100% of replicates, CRPS selects *evolved/varied* in 99%, the two rules differ in 99.2%, and S_0 ’s candidate is less dispersed by 0.0453 [0.0439, 0.0465]. That is the phenomenon Prop. 1 predicts, produced by an ordinary prompt search rather than by a constructed counterexample.

What these frequencies are, and are not. The 90% stability line that separates “stable” from “inconclusive” above is a descriptive cut we chose, not an inferential guarantee. It is not load-bearing: gpt-oss alone qualifies at 80, 90, 95 and 99%, and neither disagreement qualifies at 99.5%. A bootstrap selection frequency is the stability of a choice under resampling, not the probability that a claim is true. The resampling is stratified to match the design — 48 respondents drawn within each of four clusters, as H21 draws them — but re-selecting

inside each replicate measures selection instability on one evaluation panel, which is not an independent deployment test. We therefore read App. M as an exploratory finite-candidate result.

The cost of following it is small. Magnitude matters for the practical claim. On gpt-oss-20b the contested pair is *evolved/varied* ($S_0 = 0.7173$, CRPS = 0.1667, $\bar{\sigma} = 0.899$) against *evolved/calibrated* ($S_0 = 0.7234$, CRPS = 0.1677, $\bar{\sigma} = 0.854$). Following S_0 costs 0.0010 of CRPS — about 0.6% of the better candidate’s score — and buys a 5% reduction in dispersion. So the degenerate direction is reachable here, but the observed cost of taking it is small. We deliberately do not compare that figure to the 0.98 regret reported for model selection in App. E: that regret is normalised within a candidate set and averaged over evaluators, and this is a raw CRPS difference, so the two are not on one scale. What the two experiments jointly show is that the disagreement occurs in both settings, not that one is quantifiably more dangerous than the other.

The directives that ask for confidence do not deliver it. Across all four models neither *decisive* nor *extreme* is selected by either rule, and both *lower* S_0 relative to the unmodified prompt. The reason is not that they concentrate in the wrong place: measured against the matching base prompt they *raise* mean predictive SD in 13 of 16 model-genotype comparisons. Their names describe an intention, and the measured effect usually runs the other way, so we report them as instruments that failed rather than as evidence about where S_0 ’s optimum lies. *Calibrated*, which S_0 selects on gpt-oss, asks for the spread of a single real respondent rather than for confidence. It is not the only condition that lowered dispersion — *evolved/varied* sits below *evolved/base* on all four models, and both rules select it on Granite-4.2 and Gemma-4. A search that reliably produced low-dispersion candidates might well find more disagreement than we do; our directive set is a weak probe of that space.

N Finite-sample behaviour of the anchors

Both anchors are estimated from $n=40$ per side, which for a classifier over 120 features is a small-sample regime. We sweep n on human-only data, where the true answers are known: two disjoint samples from the same cluster must be indistinguishable, so the honest C2ST value is exactly 0.5 and any excess is estimator bias.

n	20	40	80	160	320	640
C2ST human	—	0.574	0.542	0.536	0.527	0.517
DF human	0.915	0.939	0.957	0.969	0.978	0.985
C2ST const.	—	0.914	1.000	1.000	1.000	1.000

Two corrections follow, in opposite directions, and both land on their theoretical values. Fitting $y = y_\infty + c n^{-1/2}$ over the non-degenerate range gives an extrapolated C2ST human ceiling of **0.500** and an extrapolated DF human ceiling of **1.000** — exactly what two samples from the same distribution must score. So the C2ST ceiling of 0.574 at $n=40$ is *entirely* finite-sample classifier bias and should be read with its n attached, while the DF ceiling of 0.939 is *entirely* plug-in Wasserstein bias and is pessimistic. Correcting the latter widens the gap to the simulator’s 0.783 rather than narrowing it.

Below $n \approx 40$ *this* classifier stops separating even the constant from real humans (the dashed entries above are the degenerate 0.5). That floor is a property of the configuration we used, not of C2ST: relaxing the minimum leaf size from 20 to 5 restores separation at $n=20$. We therefore recommend reporting every anchor together with the n and the classifier settings it was measured at, and treating any small- n floor as something to establish per configuration rather than to inherit from us.

References

- Qingyan Guo, Rui Wang, Junliang Guo, Bei Li, Kaitao Song, Xu Tan, Guoqing Liu, Jiang Bian, and Yujiu Yang. 2024. Connecting large language models with evolutionary algorithms yields powerful prompt optimizers. In *International Conference on Learning Representations*.
- Shashank Gupta, Vaishnavi Shrivastava, Ameet Deshpande, Ashwin Kalyan, Peter Clark, Ashish Sabharwal, and Tushar Khot. 2024. Bias runs deep: Implicit reasoning biases in persona-assigned language models. In *International Conference on Learning Representations*.
- Bin Han, Deuksin Kwon, and Jonathan Gratch. 2026. Personality expression across contexts: Linguistic and behavioral variation in LLM agents. *arXiv preprint arXiv:2602.01063*.
- Guangyuan Jiang, Manjie Xu, Song-Chun Zhu, Wenjuan Han, Chi Zhang, and Yixin Zhu. 2024. Evaluating and inducing personality in pre-trained language models. In *Advances in Neural Information Processing Systems*.
- John A. Johnson. 2014. Measuring thirty facets of the five factor model with a 120-item public domain in-

ventory: Development of the ipip-neo-120. *Journal of Research in Personality*, 51:78–89.

David Lopez-Paz and Maxime Oquab. 2017. Revisiting classifier two-sample tests. In *International Conference on Learning Representations*.

Joon Sung Park, Carolyn Q. Zou, Aaron Shaw, Benjamin Mako Hill, Carrie Cai, Meredith Ringel Morris, Robb Willer, Percy Liang, and Michael S. Bernstein. 2024. Generative agent simulations of 1,000 people. In *arXiv preprint arXiv:2411.10109*.

Archiki Prasad, Peter Hase, Xiang Zhou, and Mohit Bansal. 2023. GrIPS: Gradient-free, edit-based instruction search for prompting large language models. In *Proceedings of the 17th Conference of the European Chapter of the ACL*.

Greg Serapio-García, Mustafa Safdari, Clément Crepy, Luning Sun, Stephen Fitz, Peter Romero, Marwa Abdulhai, Aleksandra Faust, and Maja Matarić. 2023. Personality traits in large language models. *arXiv preprint arXiv:2307.00184*.

Yongchao Zhou, Andrei Ioan Muresanu, Ziwen Han, Keiran Paster, Silviu Pitis, Harris Chan, and Jimmy Ba. 2022. Large language models are human-level prompt engineers. In *International Conference on Learning Representations*.